



INDIAN INSTITUTE OF  
INFORMATION  
TECHNOLOGY

# State Space Search

Dr. Animesh Chaturvedi

Assistant Professor: IIIT Dharwad

Young Researcher: Heidelberg Laureate Forum

Postdoc: King's College London & The Alan Turing Institute

PhD: IIT Indore MTech: IIITDM Jabalpur



Indian Institute of Technology Indore  
भारतीय प्रौद्योगिकी संस्थान इंदौर



PDPM


Indian Institute of Information Technology,  
Design and Manufacturing, Jabalpur



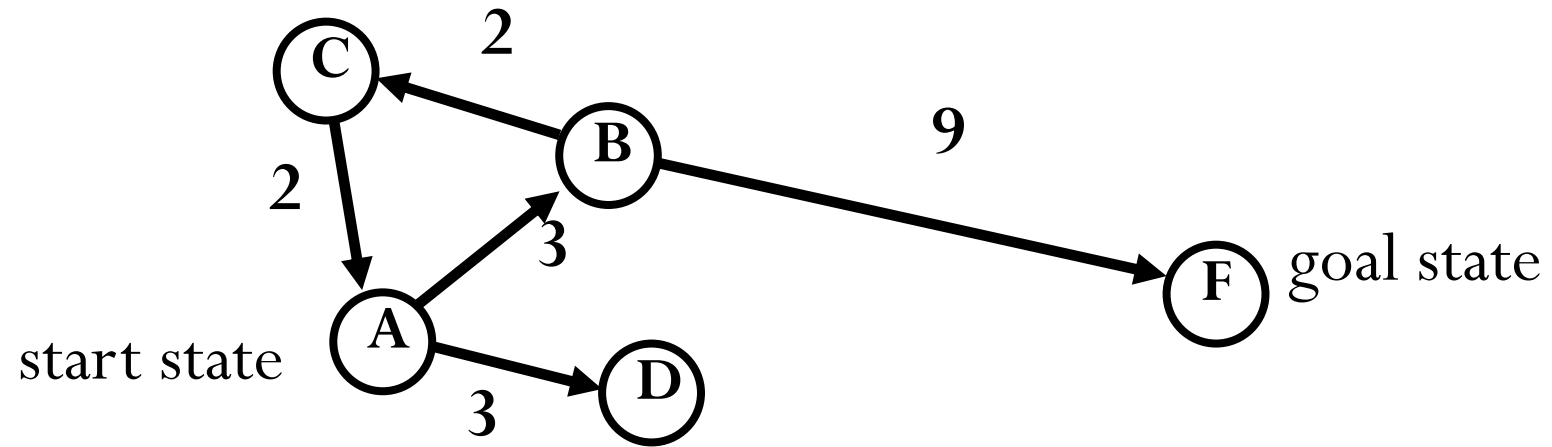
# Search

- We have some actions that can change the state of the world
  - Change induced by an action perfectly predictable
- Try to come up with a sequence of actions that will lead us to a goal state
  - May want to minimize number of actions
  - More generally, may want to minimize total cost of actions
- Do not need to execute actions in real life while searching for solution!
  - Everything perfectly predictable anyway

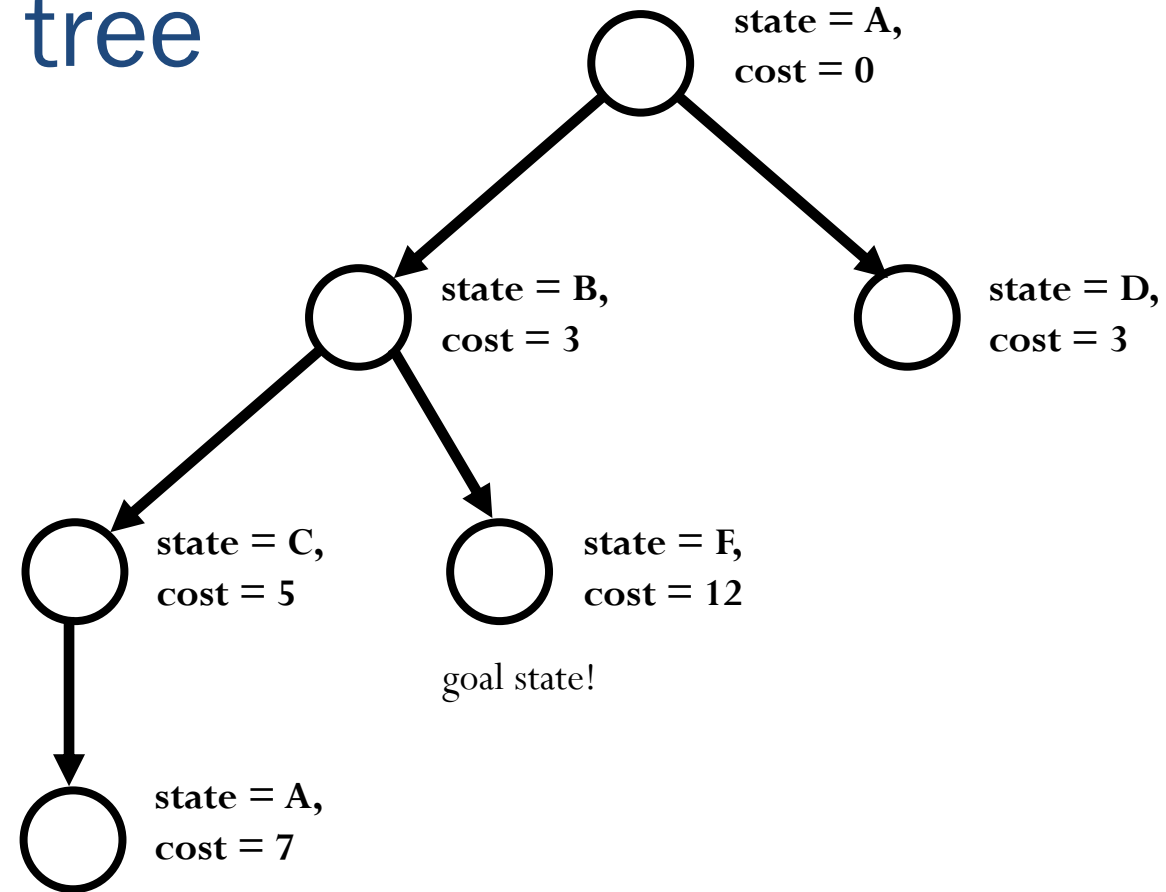
# A simple example: traveling on a graph



# Searching for a solution

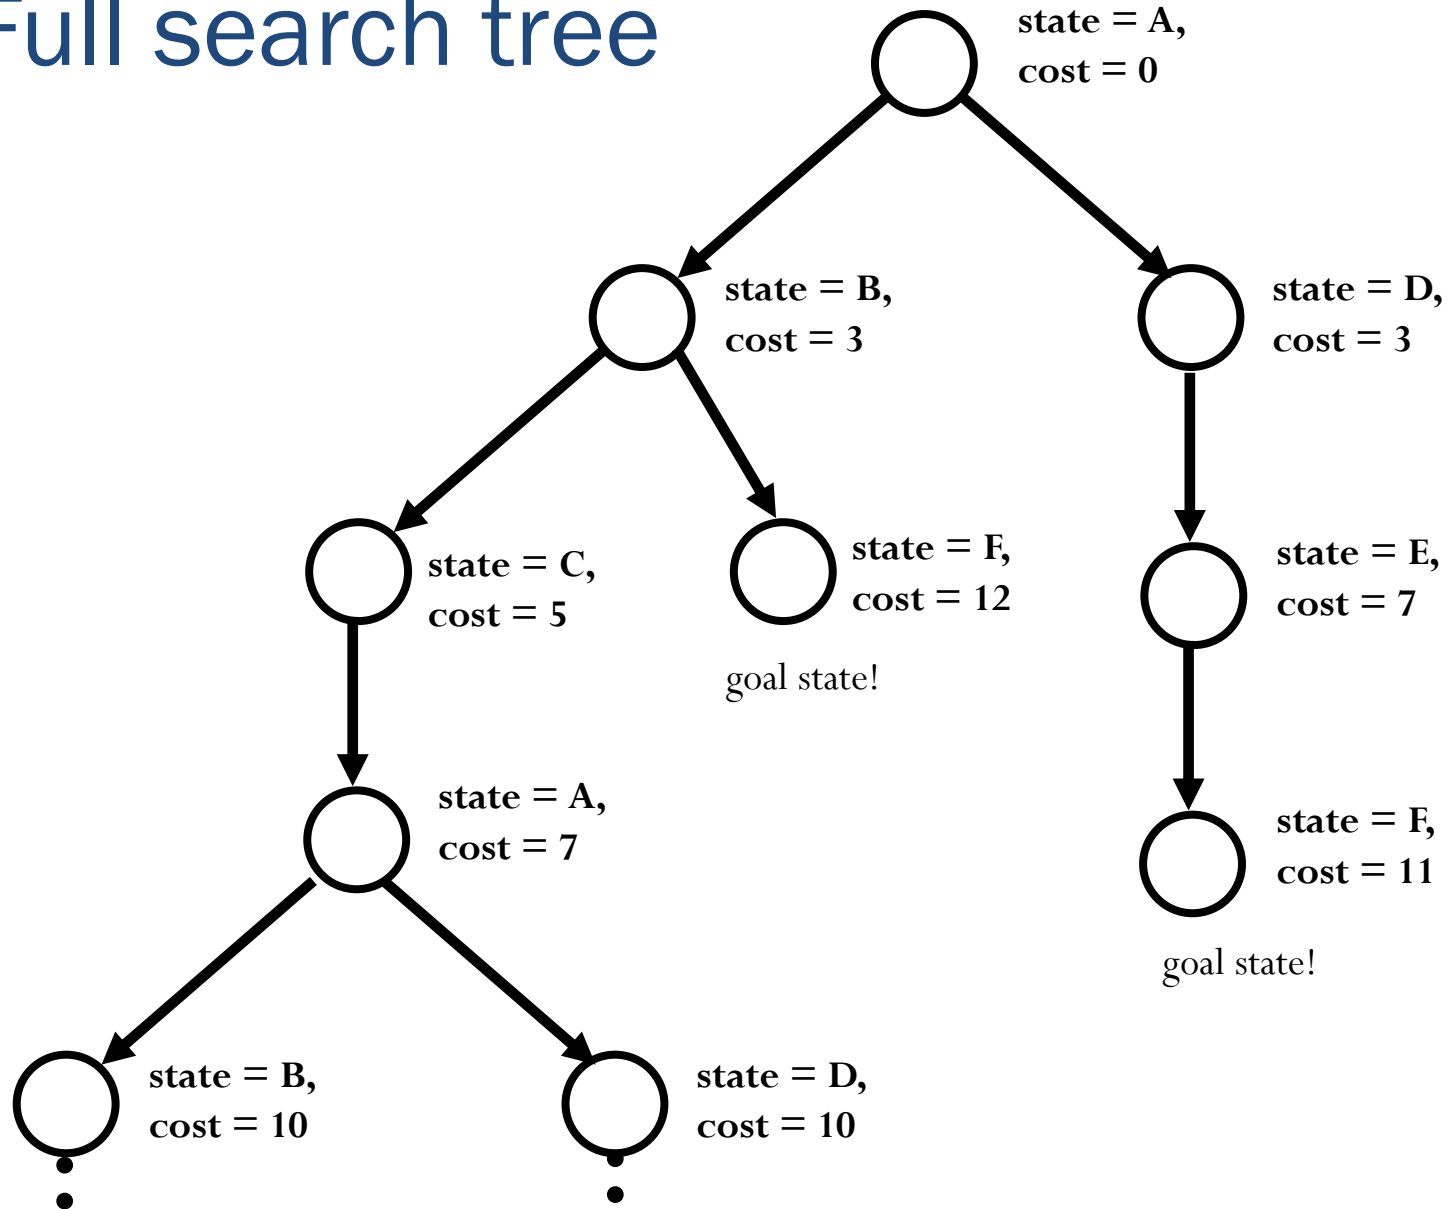


# Search tree

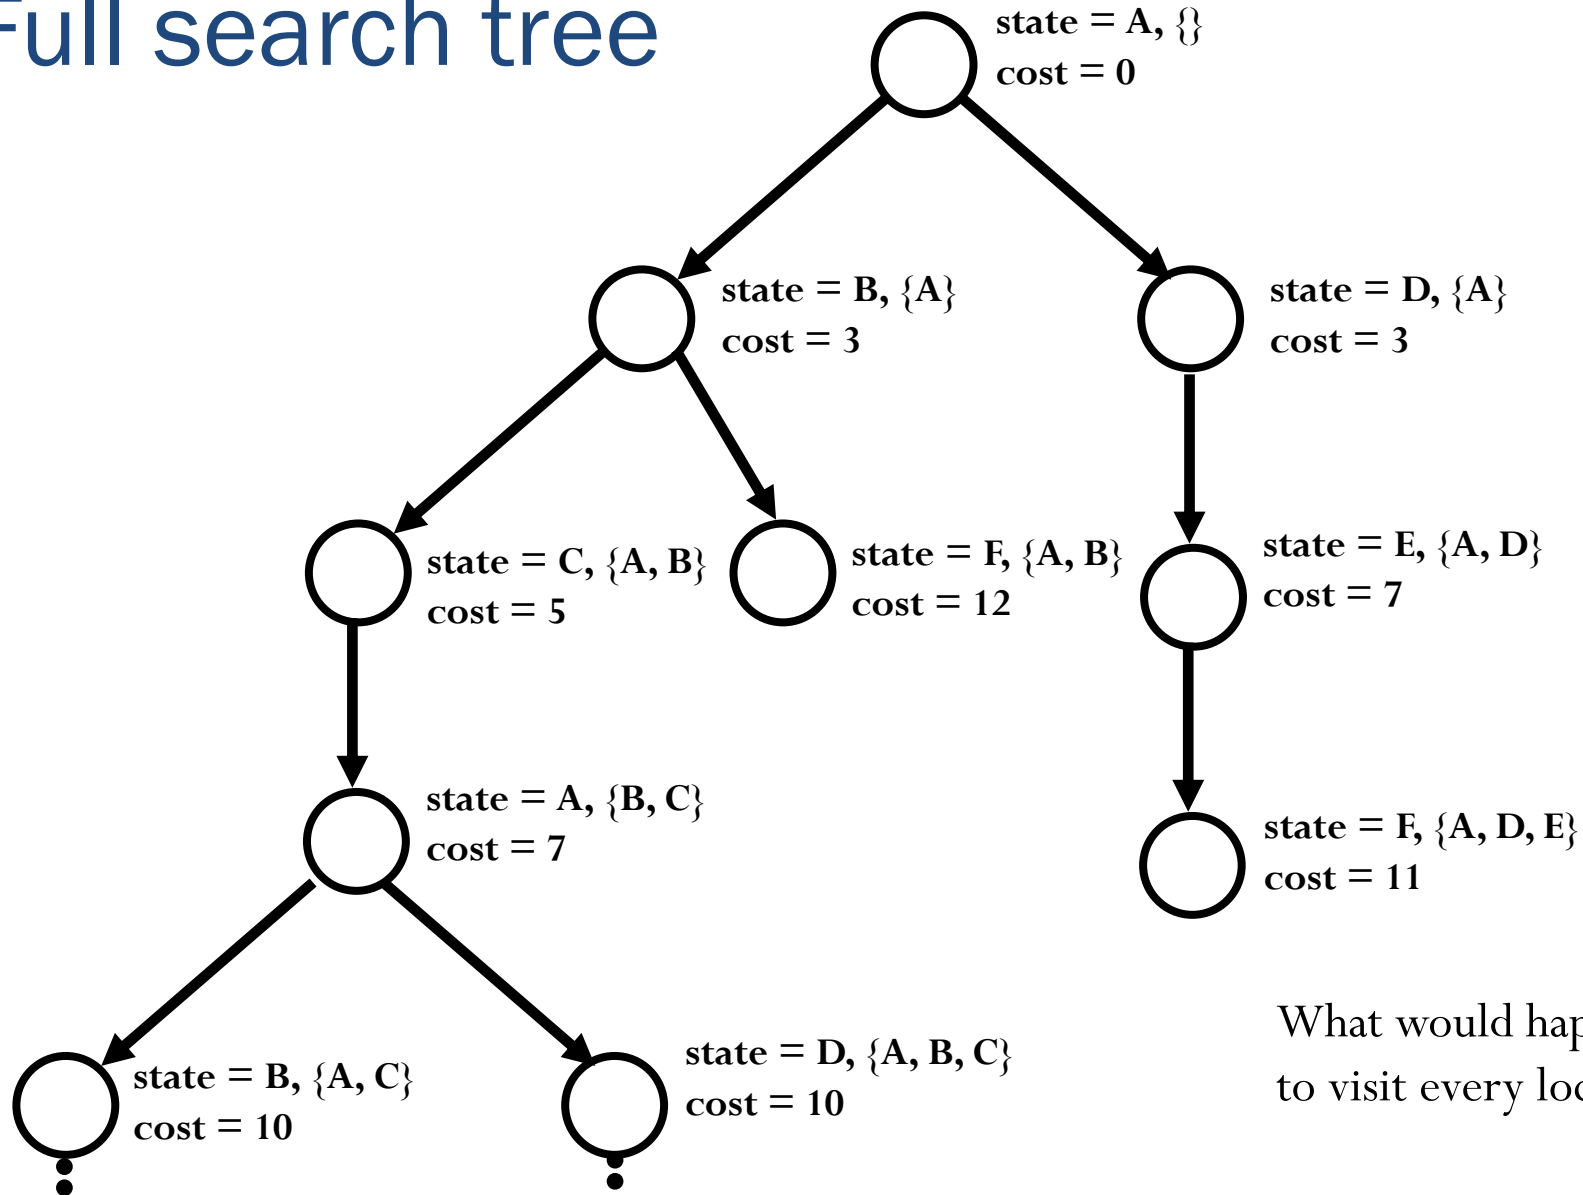


search tree nodes and states are not the same thing!

# Full search tree



# Full search tree



What would happen if the goal were to visit every location twice?

# Key concepts in search

- Set of **states** that we can be in
  - Including an **initial state**...
  - ... and **goal states** (equivalently, a **goal test**)
- For every state, a set of **actions** that we can take
  - Each action results in a new state
  - Typically defined by **successor function**
    - Given a state, produces all states that can be reached from it
- **Cost function** that determines the cost of each action (or **path** = sequence of actions)
- **Solution**: path from initial state to a goal state
  - **Optimal solution**: solution with minimal cost



# Uninformed search

# Uninformed search

- Given a state, we only know whether it is a goal state or not
- Cannot say one nongoal state looks better than another nongoal state
- Can only traverse state space blindly in hope of somehow hitting a goal state at some point
  - Also called **blind search**
  - Blind does **not** imply unsystematic!

# Searching Examples





Rush Hour: Move cars forward and backward to “escape”




8-queens problem

Logistics



8-puzzle



# Generic search algorithm

- **Fringe** = set of nodes **generated** but not **expanded**
- $\text{fringe} := \{\text{initial state}\}$
- loop:
  - if fringe empty, declare failure
  - choose and remove a node  $v$  from fringe
  - check if  $v$ 's state  $s$  is a goal state; if so, declare success
  - if not, expand  $v$ , insert resulting nodes into fringe
- Key question in search: Which of the generated nodes do we expand next?


# 8-puzzle

1		2
4	5	3
7	8	6

1	2	3
4	5	6
7	8	

goal state

# 8-puzzle




# Breadth-First Search

0

First-In-First-Out (FIFO)  
Queue


# Breadth-First Search



First-In-First-Out (FIFO)  
Queue




# Breadth-First Search




First-In-First-Out (FIFO)  
Queue

# Breadth-First Search




First-In-First-Out (FIFO)  
Queue

# Breadth-First Search




First-In-First-Out (FIFO)  
Queue

# Breadth-First Search




First-In-First-Out (FIFO)  
Queue

# Breadth-First Search




First-In-First-Out (FIFO)  
Queue

# Breadth-First Search



# Breadth-First Search




First-In-First-Out (FIFO)  
Queue

# Properties of Breadth-First Search

- Nodes are expanded in the same order in which they are generated
  - Fringe can be maintained as a First-In-First-Out (FIFO) queue
- BFS is **complete**: if a solution exists, one will be found
- BFS finds a **shallowest** solution
  - Not necessarily an optimal solution
- If every node has  $b$  successors (the **branching factor**), first solution is at depth  $d$ , then fringe size will be at least  $b^d$  at some point
  - This much space (and time) required ☹





# Depth-First Search



0


Last-In-First-Out (LIFO)  
Stack

# Depth-First Search



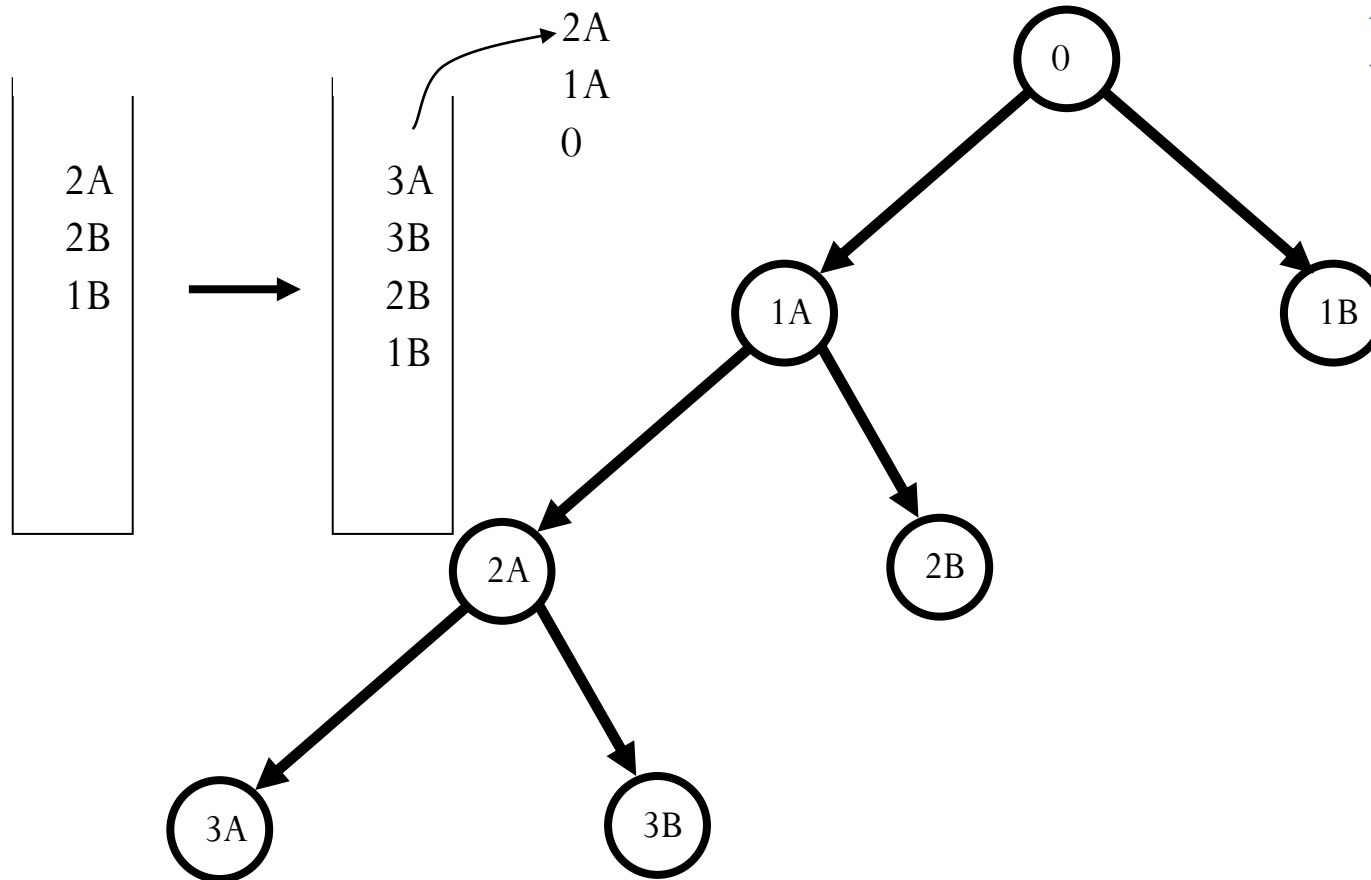
Last-In-First-Out (LIFO)  
Stack

# Depth-First Search



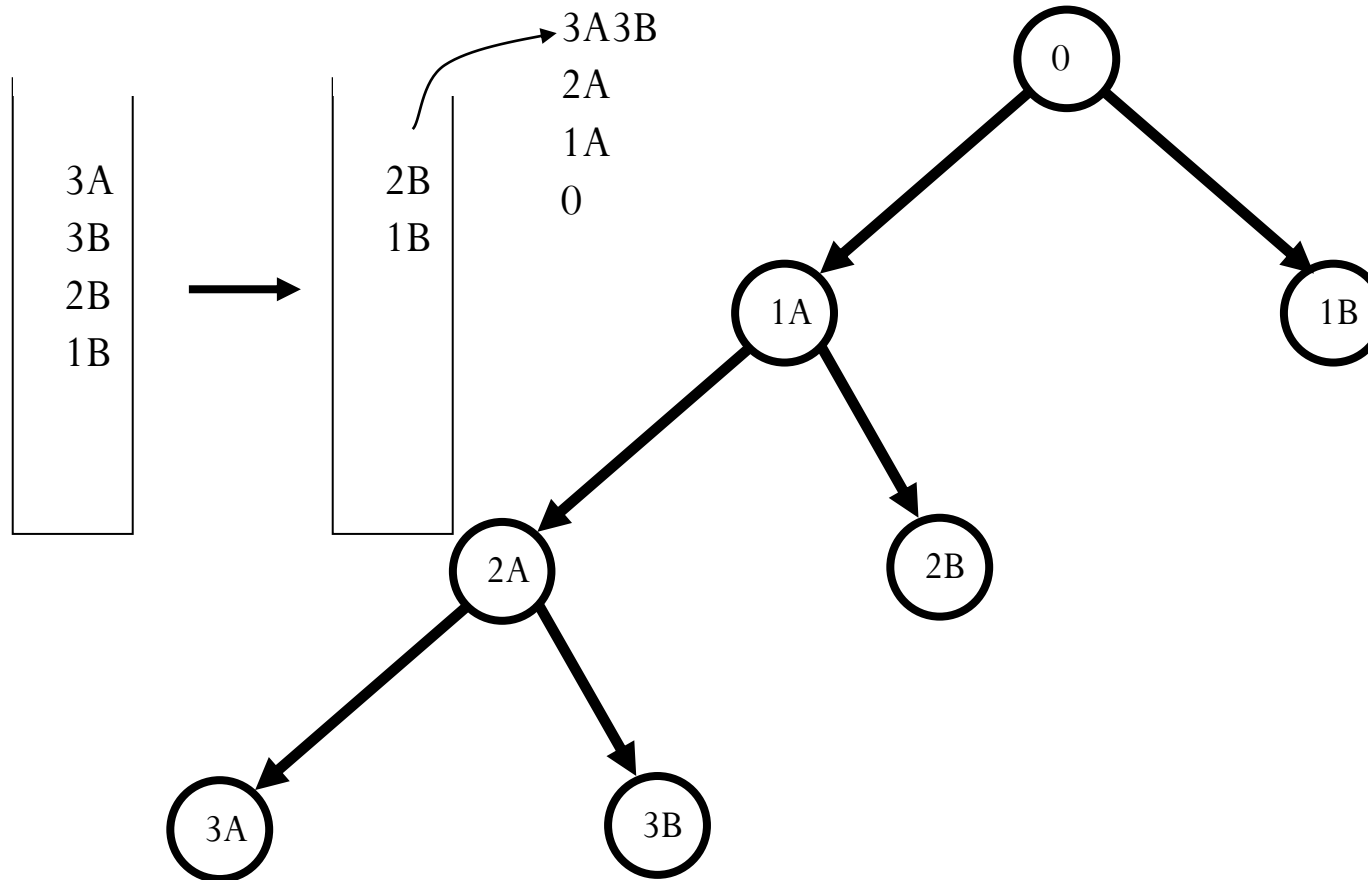
Last-In-First-Out (LIFO)  
Stack

# Depth-First Search



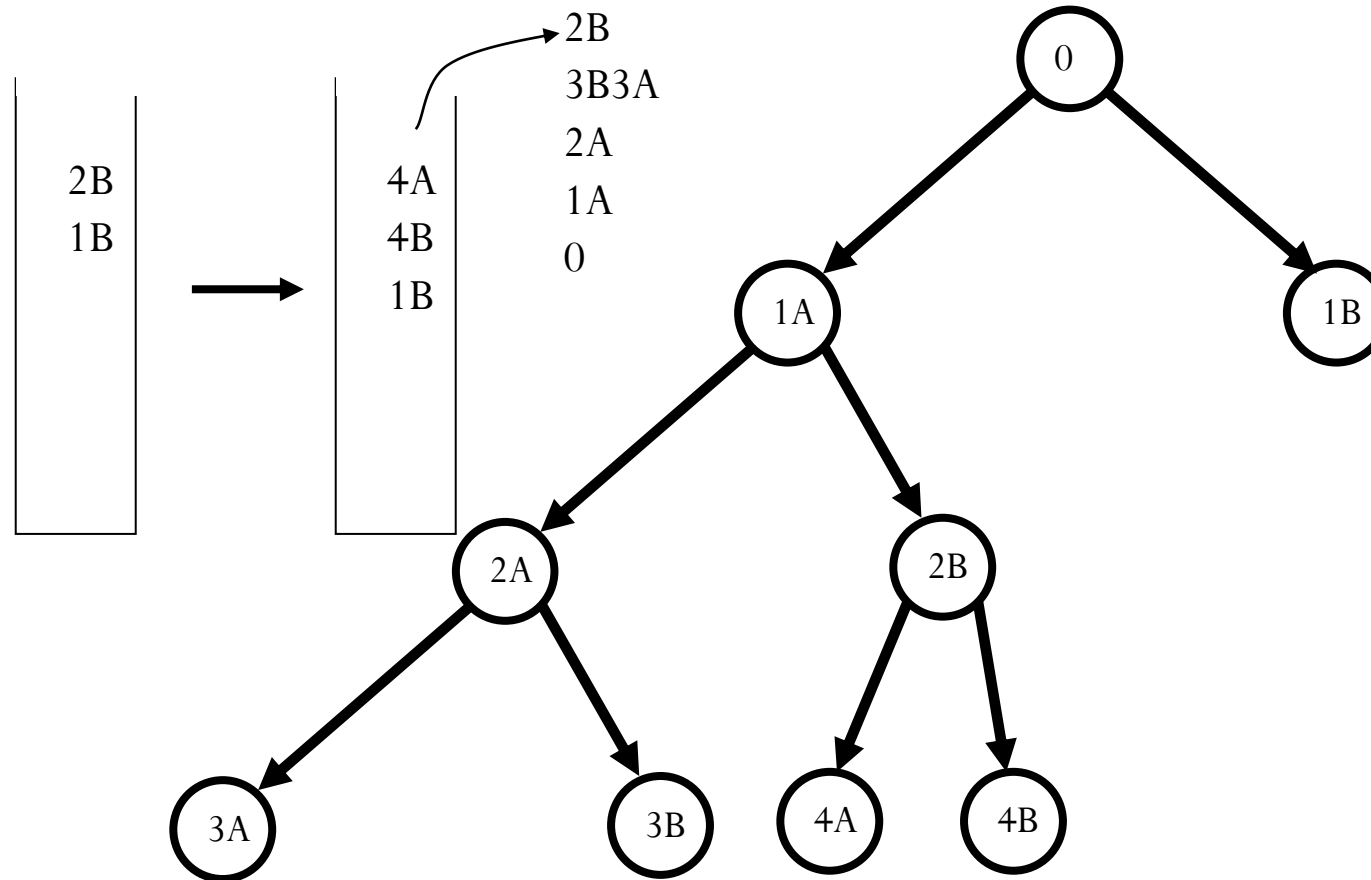
Last-In-First-Out (LIFO)  
Stack

# Depth-First Search



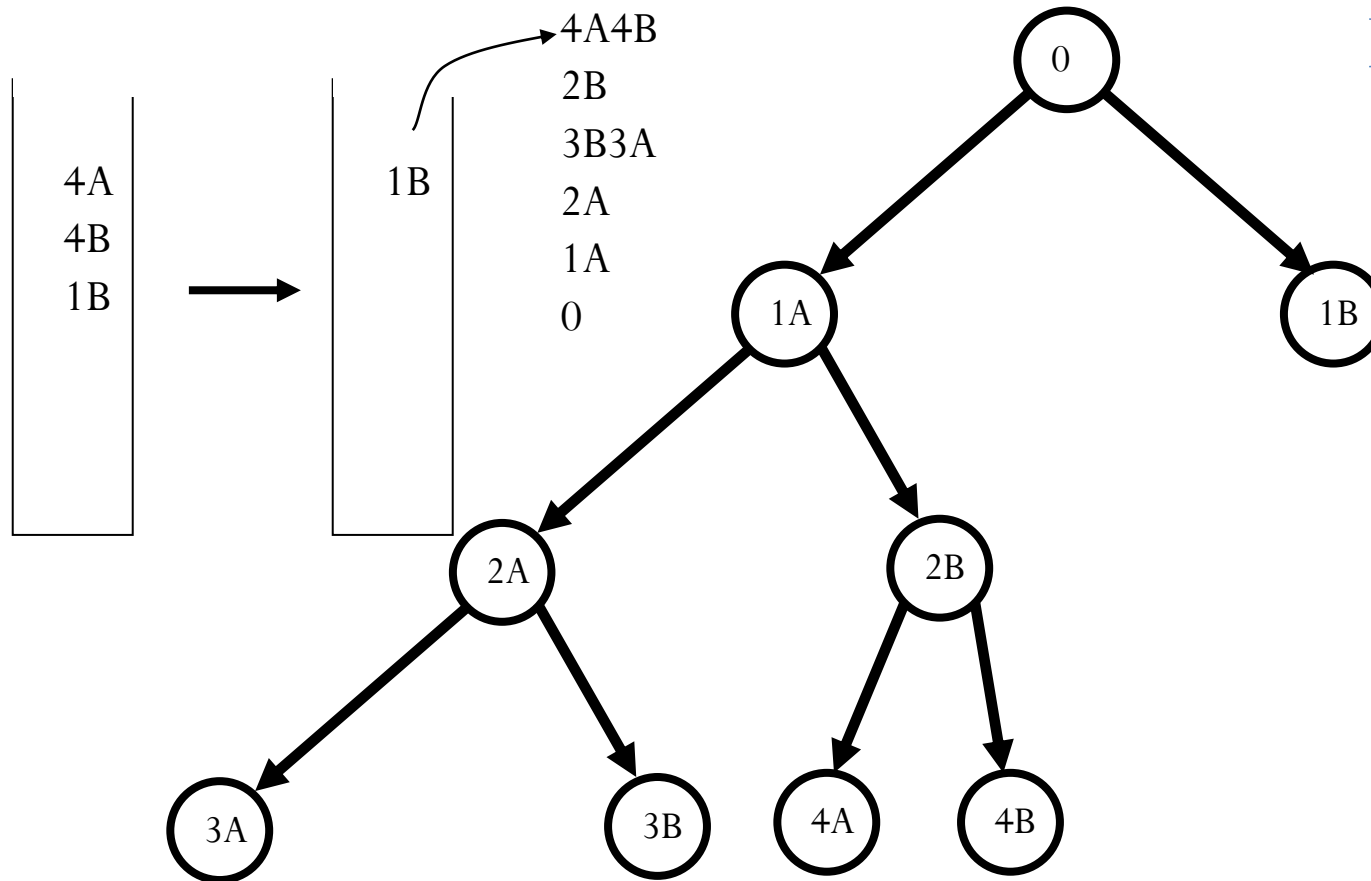
Last-In-First-Out (LIFO)  
Stack

# Depth-First Search



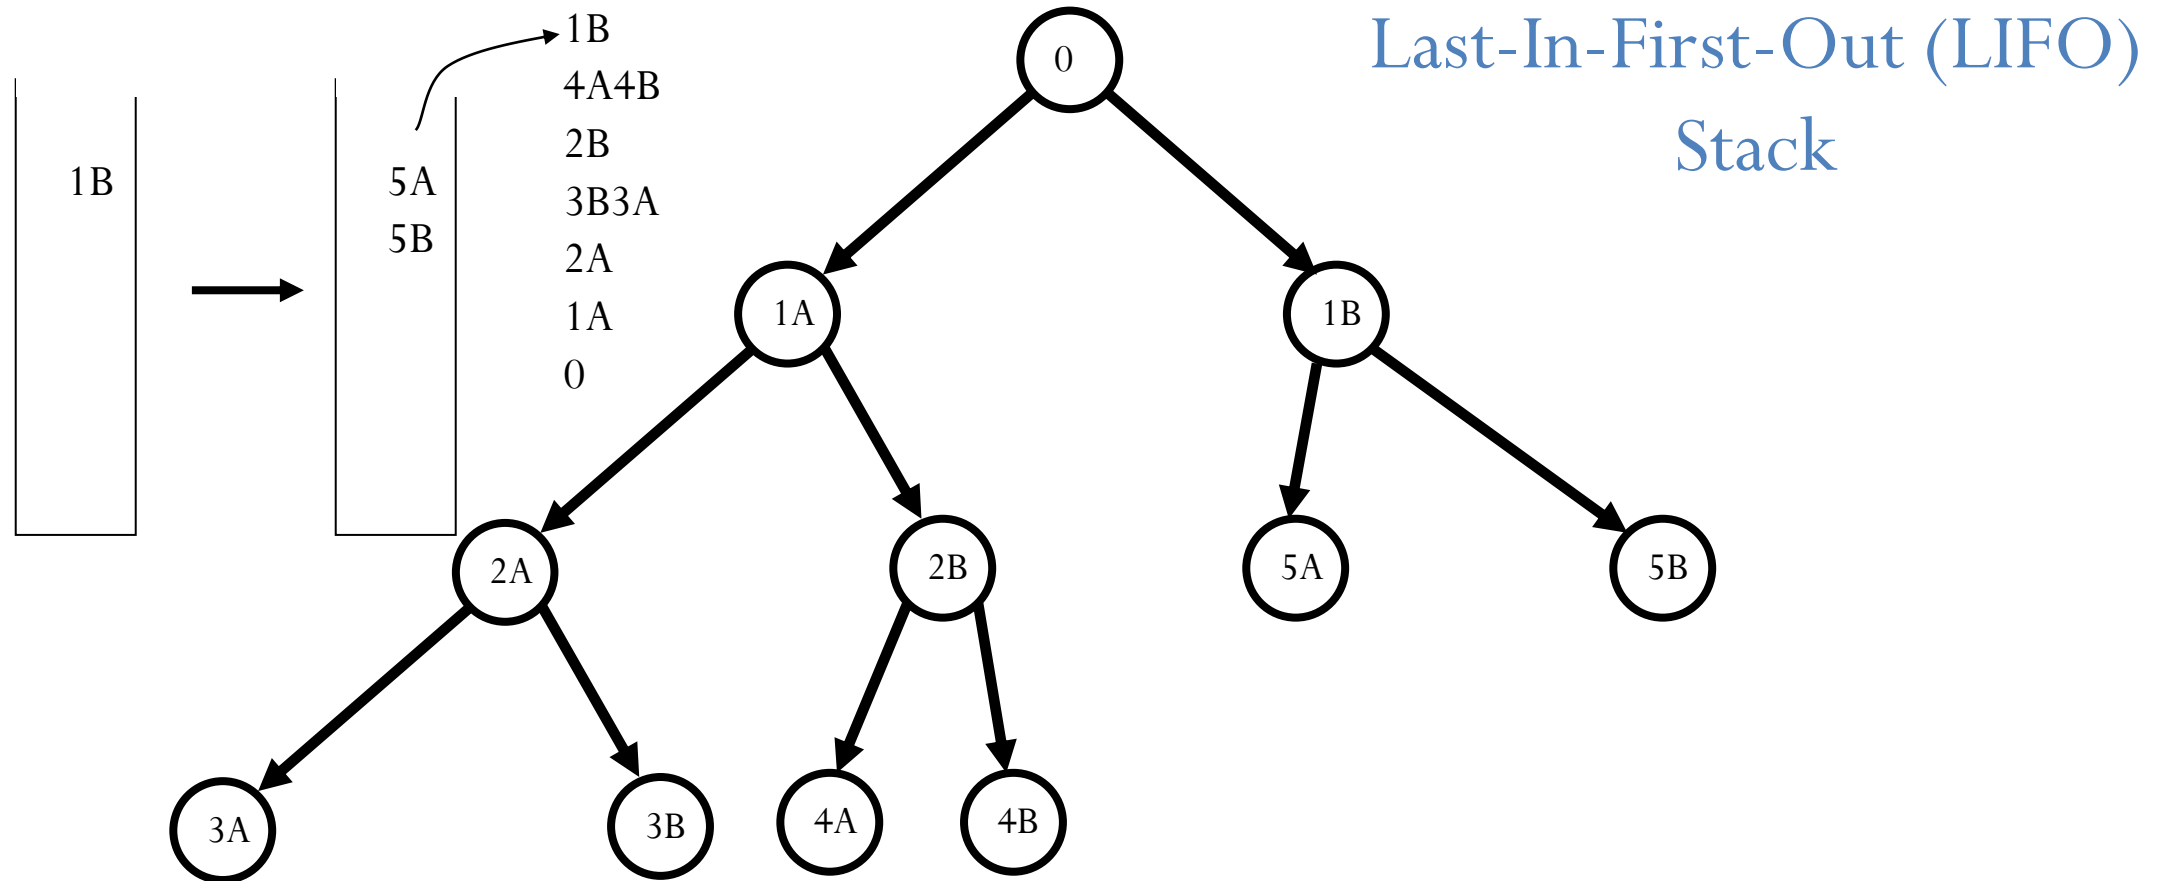
Last-In-First-Out (LIFO)  
Stack

# Depth-First Search



Last-In-First-Out (LIFO)  
Stack

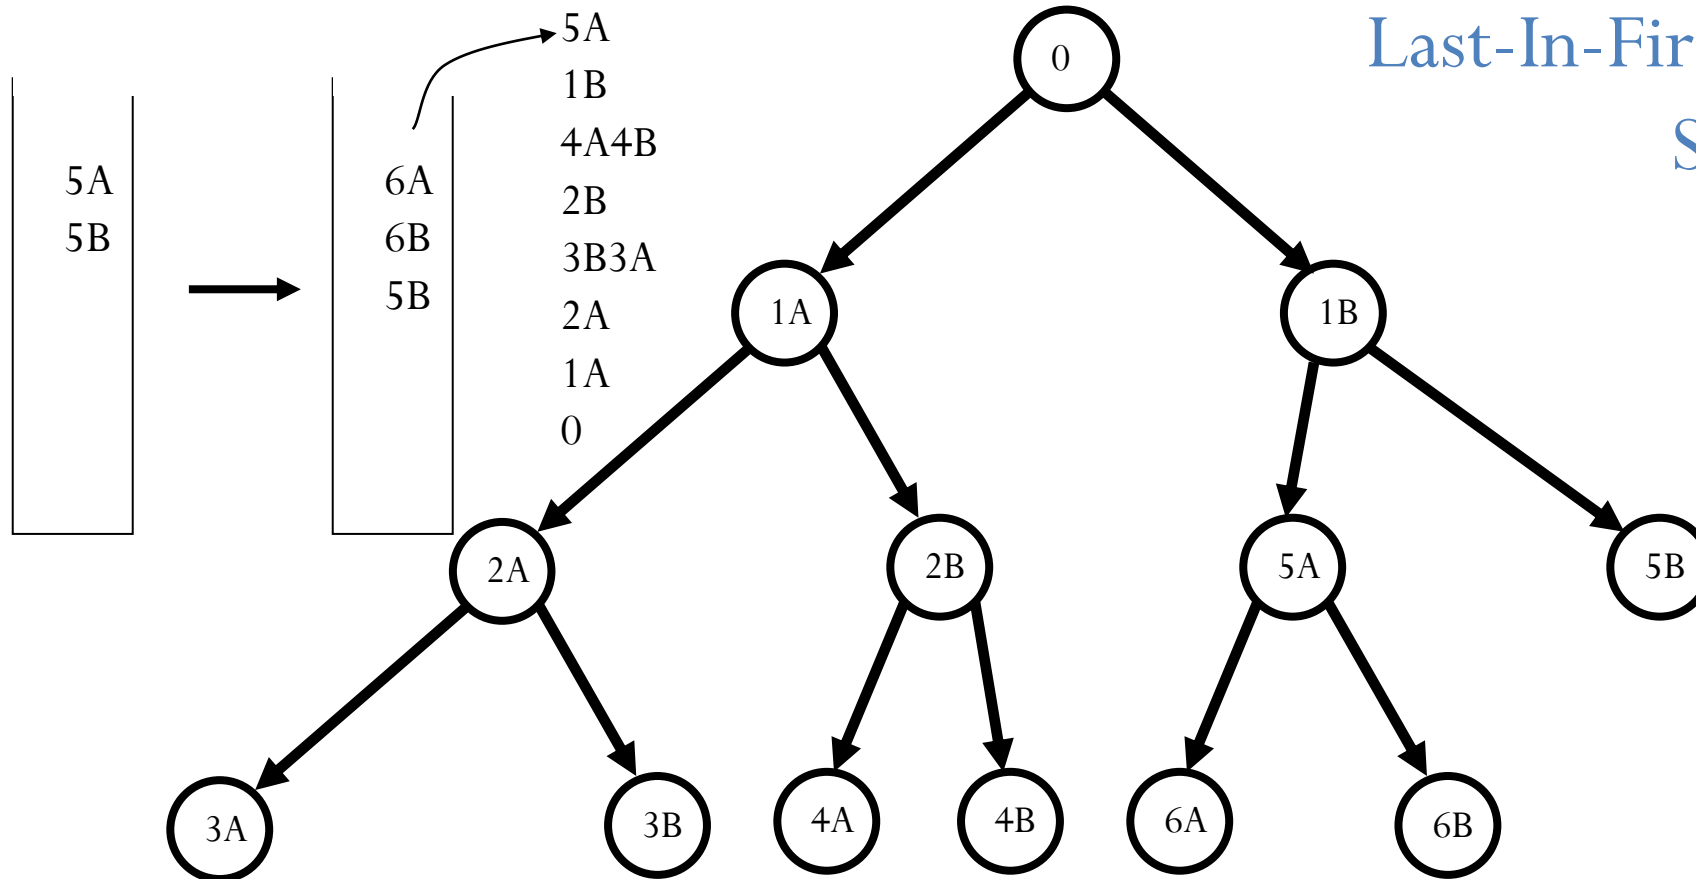
# Depth-First Search






# Depth-First Search

Last-In-First-Out (LIFO)  
Stack




# Depth-First Search



Last-In-First-Out (LIFO)  
Stack


# Depth-First Search

Last-In-First-Out (LIFO)  
Stack




# Depth-First Search

Last-In-First-Out (LIFO)  
Stack



# Depth-First Search



Last-In-First-Out (LIFO)  
Stack

# Implementing Depth-First Search

- Fringe can be maintained as a Last-In-First-Out (LIFO) queue (aka. a stack)
- Also easy to implement recursively:
- DFS(node)
  - If goal(node) return solution(node);
  - For each successor of node
    - Return DFS(successor) unless it is *failure*;
  - Return *failure*;

# Properties of depth-first search

- Not complete (might cycle through nongoal states)
- If solution found, generally not optimal/shallowest
- If every node has  $b$  successors (the **branching factor**), and we search to at most depth  $m$ , fringe is at most  $b^m$ 
  - Much better space requirement 😊
  - Actually, generally don't even need to store all of fringe
- Time: still need to look at every node
  - $b^m + b^{m-1} + \dots + 1$  (for  $b > 1$ ,  $O(b^m)$ )
  - **Inevitable** for uninformed search methods...

# Combining good properties of BFS and DFS

- **Limited depth DFS:** just like DFS, except never go deeper than some depth  $d$
- **Iterative deepening DFS:**
  - Call limited depth DFS with depth 0;
  - If unsuccessful, call with depth 1;
  - If unsuccessful, call with depth 2;
  - Etc.
- Complete, finds shallowest solution
- Space requirements of DFS
- May seem wasteful timewise because replicating effort
  - Really not that wasteful because **almost all effort at deepest level**
  - $db + (d-1)b^2 + (d-2)b^3 + \dots + 1b^d$  is  $O(b^d)$  for  $b > 1$



# Searching solution evaluation

- Comparing multiple searching algorithm based on
  - Completeness: does it always find a solution if one exist?
  - Time complexity: How long depends on number of nodes
  - Space complexity: Memory depends on number of nodes
  - Optimality: Find shortest path (or least cost solution)?
  - Systematicity: does it visit each state at most once?

# Depth vs. Breadth-first

Let  $|T(s)| \leq b$  (branching factor), goal at depth  $d$

- How implement priority queue?
- Completeness?
- Time complexity?
- Space complexity?
- Optimality?

# Breath First Search

- Completeness?
  - Yes
- Time complexity?
  - $O(b^d)$
- Space complexity?
  - $O(b^d)$  😞
- Optimality?
  - yes

# Depth First Search

- Completeness?
  - Yes, assuming state space finite
- Time complexity?
  - $O(|V+E|)$ , can do well if lots of goals
- Space complexity?
  - $O(n)$ , n deepest point of search
- Optimality?
  - No 😞

Let  $b$  as branching factor with goal at depth  $d$

# Depth-limited Search

DFS, only expand nodes depth  $\leq l$ .

- Completeness?
  - No, if  $l \leq d$ . 😞
- Time complexity?
  - $O(b^l)$
- Space complexity?
  - $O(l)$
- Optimality?
  - No 😞

# Iterative Deepening

Depth limited, increasing  $l$ .

- Completeness?
  - Yes. 😊
- Time complexity?
  - $O(b^d)$ , even with repeated work! 😊
- Space complexity?
  - $O(d)$  😊
- Optimality?
  - Yes 😊

# Bidirectional search

- Even better: search from both the start and the goal, in parallel!
- If the shallowest solution has depth  $d$  and branching factor is  $b$  on both sides, requires only  $O(b^{d/2})$  nodes to be explored!




image from [cs-alb-pc3.massey.ac.nz/notes/59302/fig03.17.gif](http://cs-alb-pc3.massey.ac.nz/notes/59302/fig03.17.gif)

# Bidirectional Search

BFS in both directions

Need  $N^{-1}$

How could this help?

- $b^l$  vs  $2b^{1/2}$

What makes this hard to implement?

# Informed Search


# Informed search

- So far, have assumed that no nonggoal state looks better than another
- Unrealistic
  - Even without knowing the road structure, some locations seem closer to the goal than others
  - Some states of the 8s puzzle seem closer to the goal than others
- Makes sense to expand closer-seeming nodes first




# Heuristics

- Key notion: **heuristic function**  $h(n)$  gives an estimate of the distance from  $n$  to the goal
  - $h(n)=0$  for goal nodes
- E.g. **straight-line distance** for traveling problem
- Say:  $h(A) = 9$ ,  $h(B) = 8$ ,  $h(C) = 9$ ,  $h(D) = 6$ ,  $h(E) = 3$ ,  $h(F) = 0$
- We're adding something new to the problem!
- Can use heuristic to decide which nodes to expand first




# Greedy best-first search

- Greedy best-first search: expand nodes with lowest  $h$  values first
- Rapidly finds the optimal solution!
- Does it always?



# A bad example for greedy

- Say:  $h(A) = 9$ ,  $h(B) = 5$ ,  $h(D) = 6$ ,  $h(E) = 3$ ,  $h(F) = 0$
- Problem: greedy evaluates the promise of a node only by how far is left to go, does not take cost occurred already into account



# IDA\*

## Memory Bounded

Just as iterative deepening gives a more memory efficient version of BFS, can define

IDA\* as a more memory efficient version of A\*.

Just use DFS with a cutoff on  $f$  values. Repeat with larger cutoff until solution found.


## What to Learn?

The A\* algorithm: its definition and behavior (finds optimal).

How to create admissible heuristics via relaxation.

# A\*

- Let  $g(n)$  be cost incurred already on path to  $n$
- Expand nodes with lowest  $g(n) + h(n)$  first
- Say:  $h(A) = 9$ ,  $h(B) = 5$ ,  $h(D) = 6$ ,  $h(E) = 3$ ,  $h(F) = 0$
- Note: if  $h=0$  everywhere, then just uniform cost search






# Admissibility

- A heuristic is **admissible**
  - if it never overestimates the distance to get the goal
  - If  $n$  is the optimal solution reachable from  $n'$ , then  $g(n) \geq g(n') + h(n')$
- Straight-line distance is admissible: can't hope for anything better than a straight road to the goal
- Admissible heuristic means that  $A^*$  is always optimistic
- $A^*$  is guaranteed to return a least-cost path from start to goal.

# Alpha Go Searching Example


# 2016: Alpha Go

- Make territories
- Capture opponent pieces
- Largest territory





# 2016: Alpha Go




# 2016: Alpha Go

- <https://www.buzzfeednews.com/article/alexkantrowitz/were-in-an-artificial-intelligence-hype-cycle>

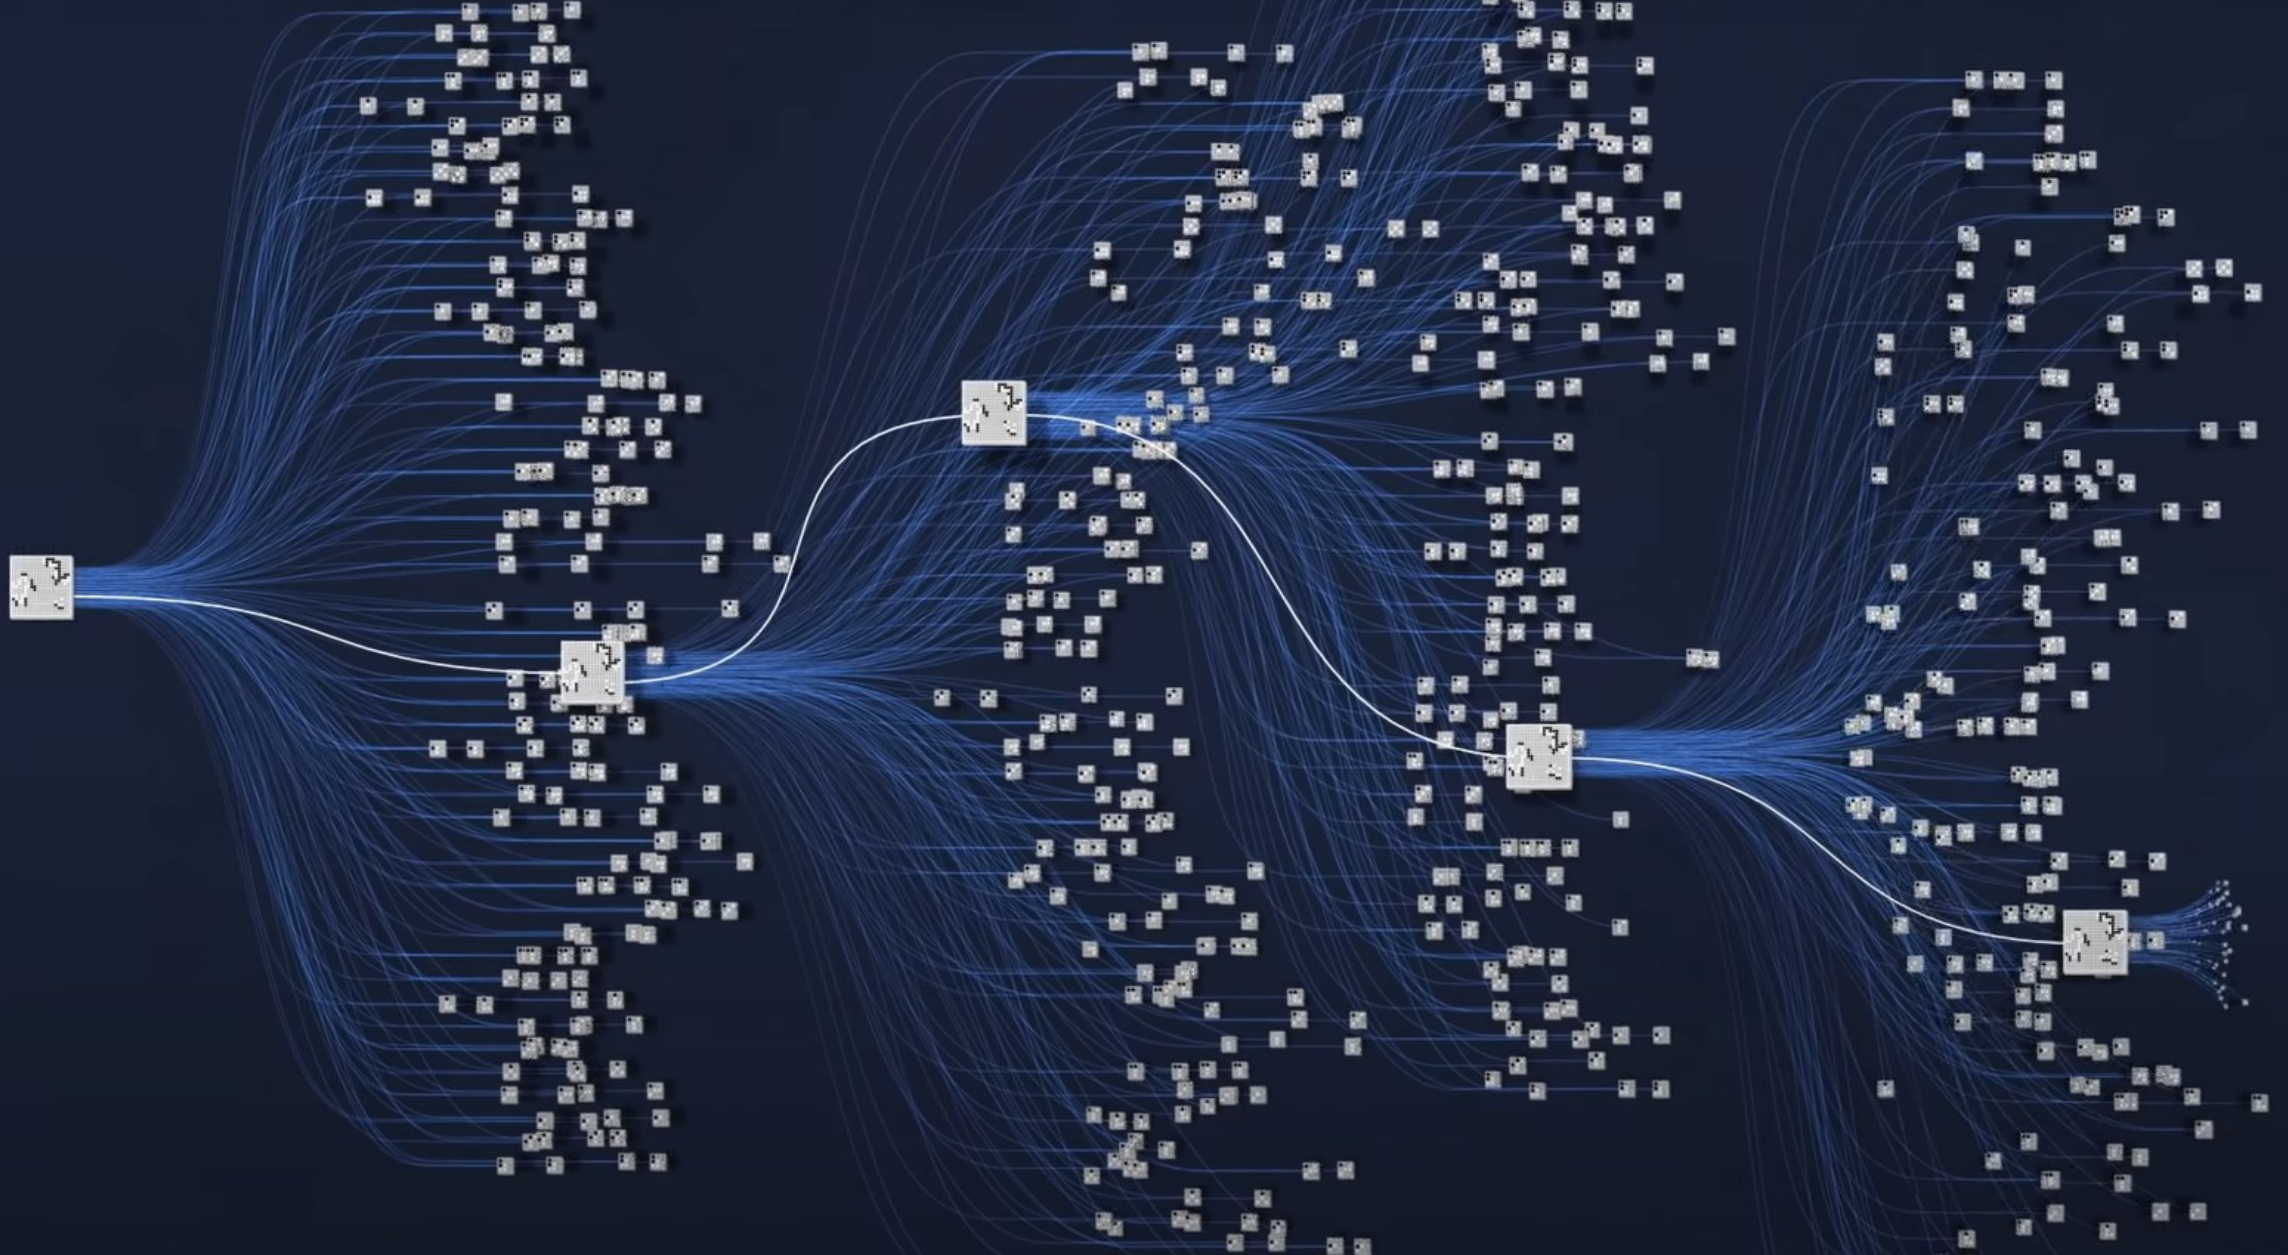


# Chess (Search Tree)





# Go (Search Tree)



# Chess v/s Go

- At the opening move in Chess there are 20 possible moves.
- Go is simpler than Chess and yet more complex.
- In large games  $O(b^d)$ , such as
  - chess ( $b \approx 35$ ,  $d \approx 80$ )
  - Go ( $b \approx 250$ ,  $d \approx 150$ ),
- Exhaustive search is infeasible, but the effective search space can be reduced by two general principles.
  - the depth of the search may be reduced by position evaluation: truncating (cut) the search tree at state  $s$  and replacing the subtree below  $s$  by an approximate value function  $v(s) \approx v^*(s)$  that predicts the outcome from state  $s$ .
  - the breadth of the search may be reduced by sampling actions from a policy  $p(a | s)$  that is a probability distribution over possible moves  $a$  in position  $s$ .

# References

- Stuart Russel, and Peter Norvig. "Artificial intelligence: A modern approach. third edit." Upper Saddle River, New Jersey 7458 (2015).
- Introduction to Artificial Intelligence, Michael L. Littman, Fall 2001  
[mlittman@cs.brown.edu](mailto:mlittman@cs.brown.edu) <https://courses.cs.duke.edu/fall08/cps270/>
- Vincent Conitzer, Artificial Intelligence <http://www.cs.duke.edu/courses/fall08/>
- Silver, David, et al. "Mastering the game of Go with deep neural networks and tree search." Nature 529.7587 (2016): 484-489.
- AlphaGo - The Movie | Full award-winning documentary "DeepMind"  
<https://www.youtube.com/watch?v=WXuK6gekU1Y>

ขอบคุณ

Thai

Grazie  
Italian

תודה רבה  
Hebrew

धन्यवादः  
Sanskrit

ಧನ್ಯವಾದಗಳು  
Kannada

Ευχαριστώ  
Greek

Thank You  
English

Gracias  
Spanish

Спасибо  
Russian

Obrigado  
Portuguese

شكراً  
Arabic

<https://sites.google.com/site/animeshchaturvedi07>

Merci  
French

多謝  
Traditional  
Chinese

धन्यवाद  
Hindi

Danke  
German

多谢  
Simplified  
Chinese

நன்றி  
Tamil

ありがとうございました  
Japanese

감사합니다  
Korean