



INDIAN INSTITUTE OF
INFORMATION
TECHNOLOGY

Data Science, Knowledge Discovery, and Data mining

Dr. Animesh Chaturvedi

Assistant Professor: IIIT Dharwad

Young Researcher: Heidelberg Laureate Forum

Postdoc: King's College London & The Alan Turing Institute

PhD: IIT Indore MTech: IIITDM Jabalpur



Indian Institute of Technology Indore
भारतीय प्रौद्योगिकी संस्थान इंदौर



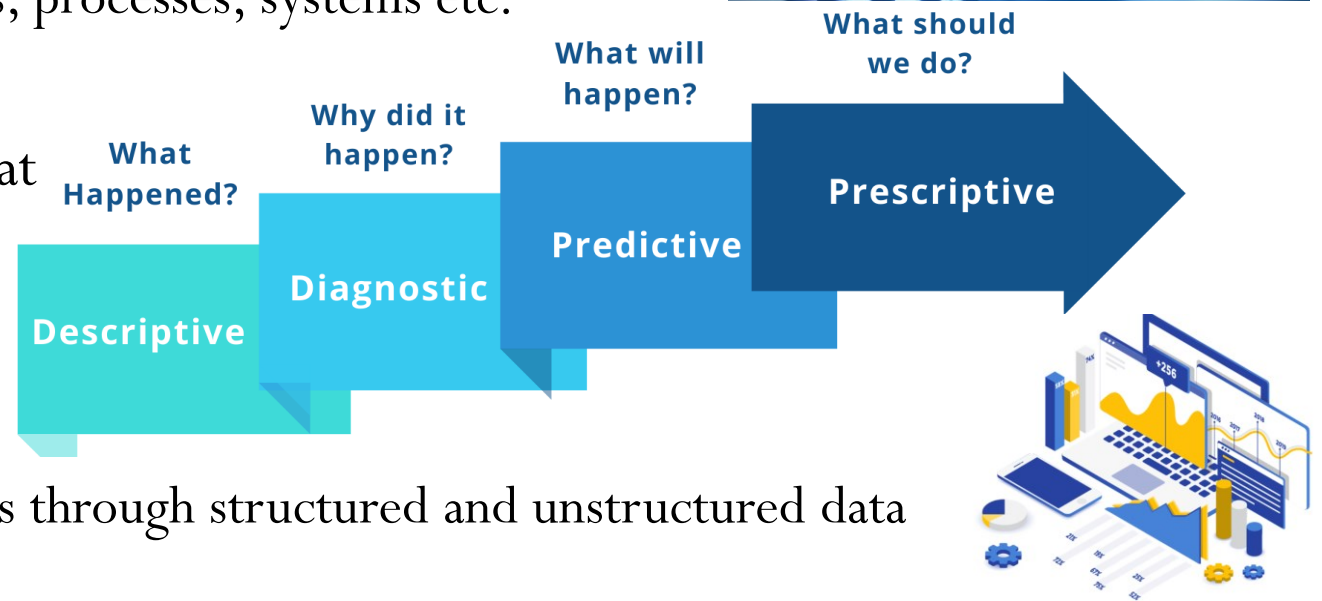
PDPM

Indian Institute of Information Technology,
Design and Manufacturing, Jabalpur



Data Science and Data Analytics

- **Data Science:** interdisciplinary science that
 - deals with data: methods, algorithms, processes, systems etc.
 - Theory oriented
- **Data Analytics:** analysis of data that
 - discovers trends, graph, tables etc.
 - Technology oriented
- Both
 - extracts knowledge and apply actions through structured and unstructured data insights
 - deals with data mining, machine learning, data management, and big data.
- **Data Scientist** and **Data Analyst** applies Data Science and Data Analytics



Evolution to Data Science Societies

- Institute of Radio Engineers (**IRE**) in 1951
- Institute of Electrical and Electronics Engineers (**IEEE**), **IEEE Computer Society** (1946; 1963)
- Association for Computing Machinery (**ACM**) in 1947
- Technology Engineering and Management Society (**TEMS**) 1955
- IEEE Systems, Man, and Cybernetics Society (**SMC**) (1958; 1972)
- Association for Computational Linguistics (**ACL**) 1962
- International Joint Conf. on Artificial Intelligence (**IJCAI**) 1969
- Special Interest Group on Management of Data (**SIG-MOD**), 1975
- Special Interest Group on Information Retrieval (**SIG-IR**) 1978
- Association for the Advancement of Artificial Intelligence (**AAAI**) 1979
- International Conf. on Machine Learning (**ICML**) 1980 in Pittsburgh
- Conf. on Neural Information Processing Systems (**NeurIPS**) 1986 -1987
- 38th IEEE International Conf. on Data Engineering (**IEEE ICDE**)*
- International World Wide Web Conf. (**WWW**) - Web Conf. 1994
- 28th International Conf. on Knowledge Discovery and Data Mining (**SIG-KDD**)*
- 22nd IEEE International Conf. on Data Mining (**IEEE ICDM**)*
- 10th IEEE International Conf. on Big Data (**IEEE Big Data**)*
- 9th IEEE International Conf. on Data Science and Advanced Analytics (**IEEE DSAA**)*

IRE

IEEE

ACM

TEMS

SMC

ACL

IJCAI

SIG-MOD & IR

AAAI, ICML

NeurIPS


ICDE, WWW

KDD, ICDM

Big Data,

DSAA

* in 2022



DBLP & Google Scholar


- “Data” in DBLP search*
 - There are 459+ Venues matched in the DBLP, world most referred computer science bibliography website.
 - DBLP launched in 1993
- “Data” in Google Scholar metrics
 - Top 20 publications venues matching “data”

*2022

	Publication*	<u>h5-index</u>	<u>h5-median</u>
1.	ACM SIGKDD International Conference on Knowledge Discovery & Data Mining	114	196
2.	IEEE Transactions on Knowledge and Data Engineering	88	147
3.	International Conference on Artificial Intelligence and Statistics	85	119
4.	ACM International Conference on Web Search and Data Mining	69	133
5.	Journal of Big Data	55	104
6.	IEEE International Conference on Data Mining	53	81
7.	IEEE International Conference on Big Data	52	93
8.	Knowledge and Information Systems	51	76
9.	ACM Conference on Recommender Systems	47	111
10.	Wiley Interdisciplinary Reviews: Data Mining and Knowledge Discovery	47	89
11.	European Conference on Machine Learning and Knowledge Discovery in Databases	40	57
12.	ACM Transactions on Intelligent Systems and Technology (TIST)	38	72
13.	Data Mining and Knowledge Discovery	37	72
14.	SIAM International Conference on Data Mining (SDM)	35	60
15.	ACM Transactions on Knowledge Discovery from Data (TKDD)	35	53
16.	International Conference on Advances in Social Networks Analysis and Mining	34	65
17.	Big Data Mining and Analytics	31	39
18.	International Journal of Data Science and Analytics	30	52
19.	Social Network Analysis and Mining	30	46
20.	Pacific-Asia Conference on Knowledge Discovery and Data Mining (PAKDD)	29	45

CSRankings.org for India

CSRankings: Computer Science Rankings

CSRankings is a metrics-based ranking of top computer science institutions around the world. Click on a triangle (▶) to expand areas or institutions. Click on a name to go to a faculty member's home page. Click on a chart icon (the  after a name or institution) to see the distribution of their publication areas as a . Click on a Google Scholar icon (🔍) to see publications, and click on the DBLP logo (📄) to go to a DBLP entry. Applying to grad school? Read this first. Do you find CSRankings useful? Sponsor CSRankings on GitHub.

Rank institutions in by publications from to

All Areas ☐ off | ☐ on

AI ☐ off | ☐ on

- ▶ Artificial intelligence ☒
- ▶ Computer vision ☒
- ▼ Machine learning & data mining ☒
- ACM SIGKDD, IMLS, NEURIPS/NIPS
- ICML ☒
- KDD ☒
- NeurIPS/NIPS ☒
- ▶ Natural language processing ☒
- ▶ The Web & information retrieval ☒

Systems ☐ off | ☐ on

- ▶ Computer architecture ☒
- ▶ Computer networks ☒
- ▶ Computer security ☒
- ▶ Databases ☒
- ▶ Design automation ☒
- ▶ Embedded & real-time systems ☒
- ▶ High-performance computing ☒
- ▶ Mobile computing ☒

#	Institution	Count	Faculty
1	▶ IISc Bangalore 	2.6	28
2	▶ IIT Bombay 	2.1	24
2	▶ IIT Delhi 	2.1	22
4	▶ IIT Kanpur 	2.0	19
5	▶ IIT Madras 	1.9	23
6	▶ IIIT Hyderabad 	1.6	19
6	▶ IIT Kharagpur 	1.6	21
8	▶ IIIT Delhi 	1.5	24
9	▶ IIT Gandhinagar 	1.3	9
9	▶ IIT Hyderabad 	1.3	10
11	▶ IIIT Bangalore 	1.2	5
11	▶ IIT Patna 	1.2	4
11	▶ IMSc 	1.2	5
11	▶ Tata Inst. of Fundamental Research 	1.2	9
15	▶ CMI 	1.1	8
15	▶ IIT Goa 	1.1	4

#S	Institution	Count	Faculty
1	▶ IISc Bangalore	2.6	28
2	▶ IIT Bombay	2.1	24
2	▶ IIT Delhi	2.1	22
4	▶ IIT Kanpur	2.0	19
5	▶ IIT Madras	1.9	23
6	▶ IIIT Hyderabad	1.6	19
6	▶ IIT Kharagpur	1.6	21
8	▶ IIIT Delhi	1.5	24
9	▶ IIT Gandhinagar	1.3	9
9	▶ IIT Hyderabad	1.3	10
11	▶ IIIT Bangalore	1.2	5
11	▶ IIT Patna	1.2	4
11	▶ IMSc	1.2	5
11	▶ Tata Inst. of Fundamental Research	1.2	9
15	▶ CMI	1.1	8
15	▶ IIT Goa	1.1	4
15	▶ IIT Guwahati	1.1	7
15	▶ IIT Jodhpur	1.1	5
15	▶ IIT Ropar	1.1	2
20	▶ BITS Pilani	1.0	1
20	▶ BITS Pilani-Goa	1.0	1
20	▶ DAIICT	1.0	2
20	▶ IIT (BHU) Varanasi	1.0	1
20	▶ IIT Indore	1.0	3
20	▶ IIT Mandi	1.0	1
20	▶ IIT Roorkee	1.0	3

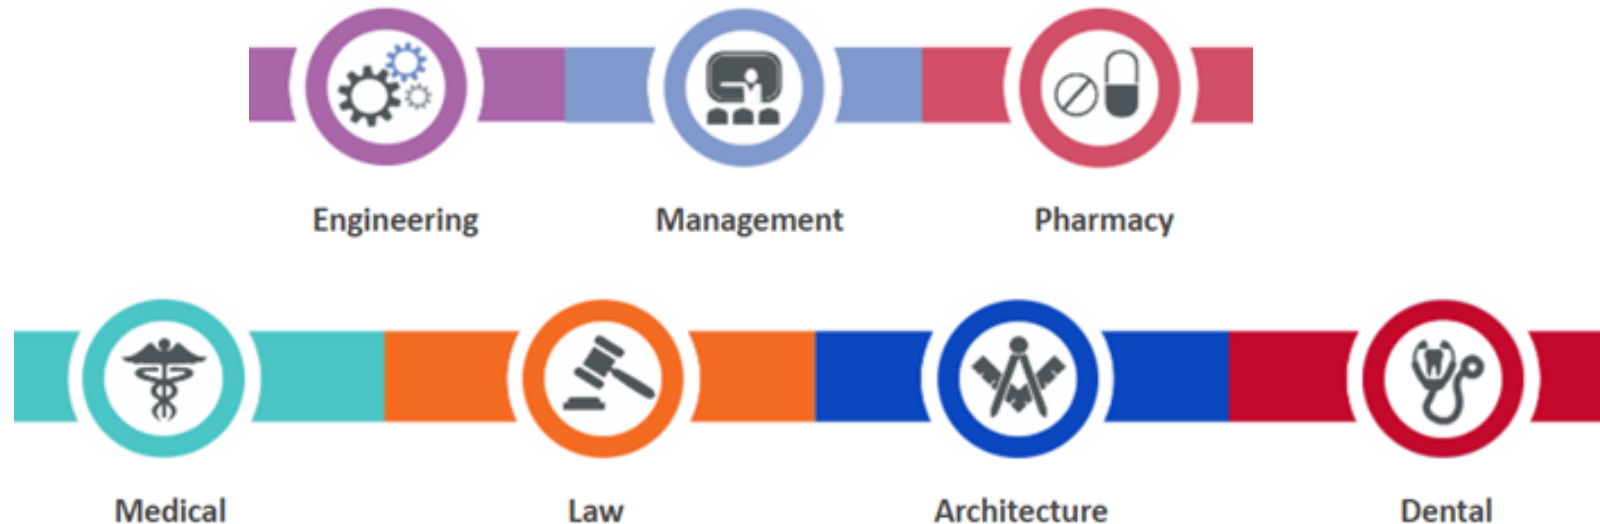
CS, EC, and IT and other related
branch based institute ranking

Institutes from India listed for A* CS events like KDD in 2022

Data Analytics for Information Technology (IT)

- IT based category parameter are different from Engineering, Maths, and Sciences
- IT means (CS-EC-DS-AI-Bio/Chem-informatics-etc)
 - research, jobs, in DBLP, Scholar Metrics
 - Competitions like ICPC, KDD Cup, ACM Student Research, etc.

IT is applied as advanced field for



Knowledge Discovery and Data Mining

- “a unifying framework for Knowledge Discovery in Database (KDD)”
- links between data mining, knowledge discovery, and other related field

Selection, Preprocessing, Transformation, Data Mining, Interpretation/Evaluation

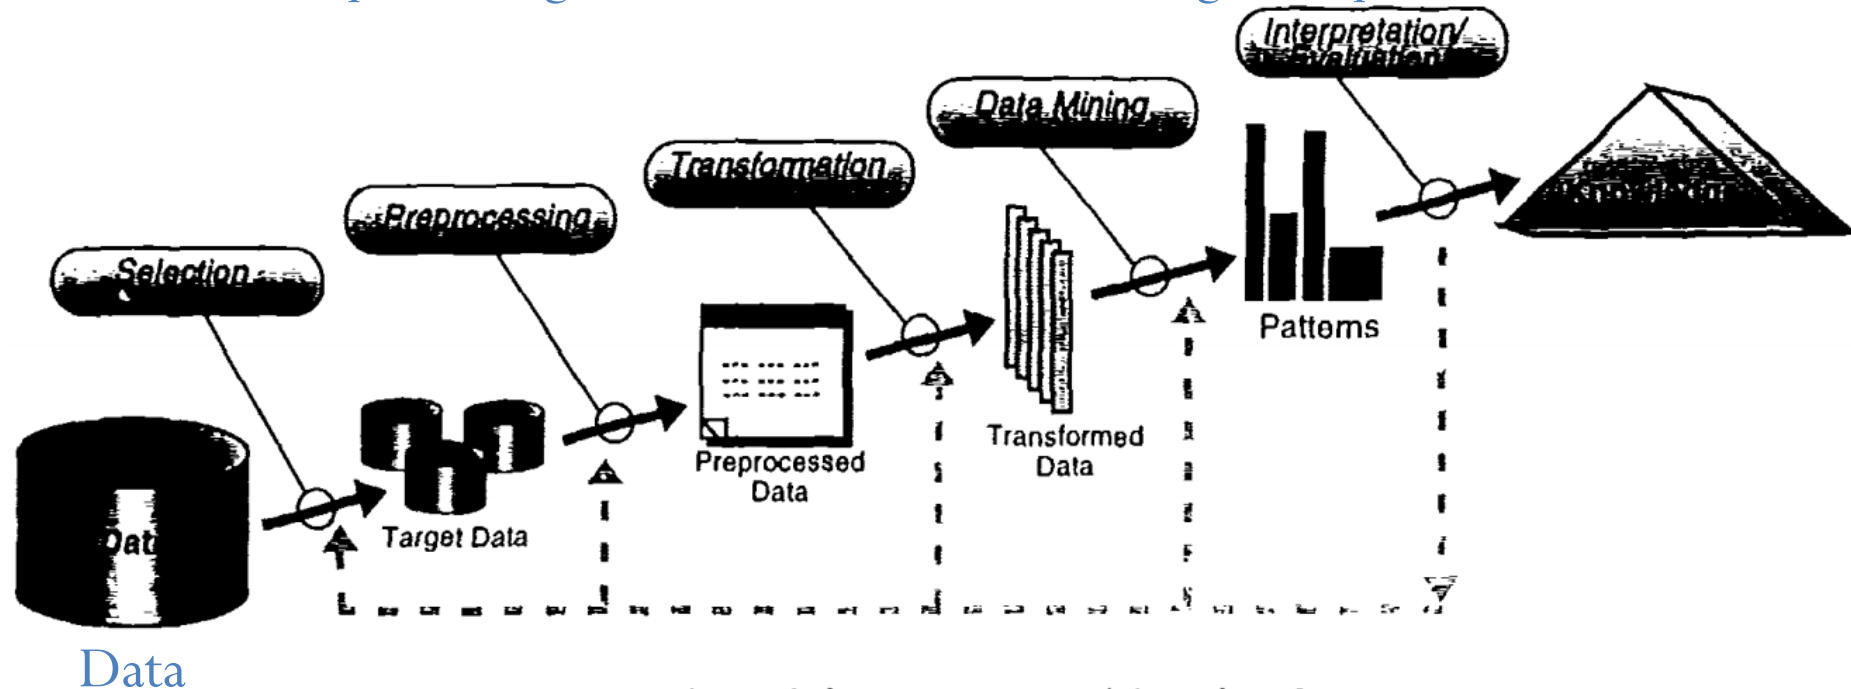


Figure 1: An overview of the steps comprising the KDD process.

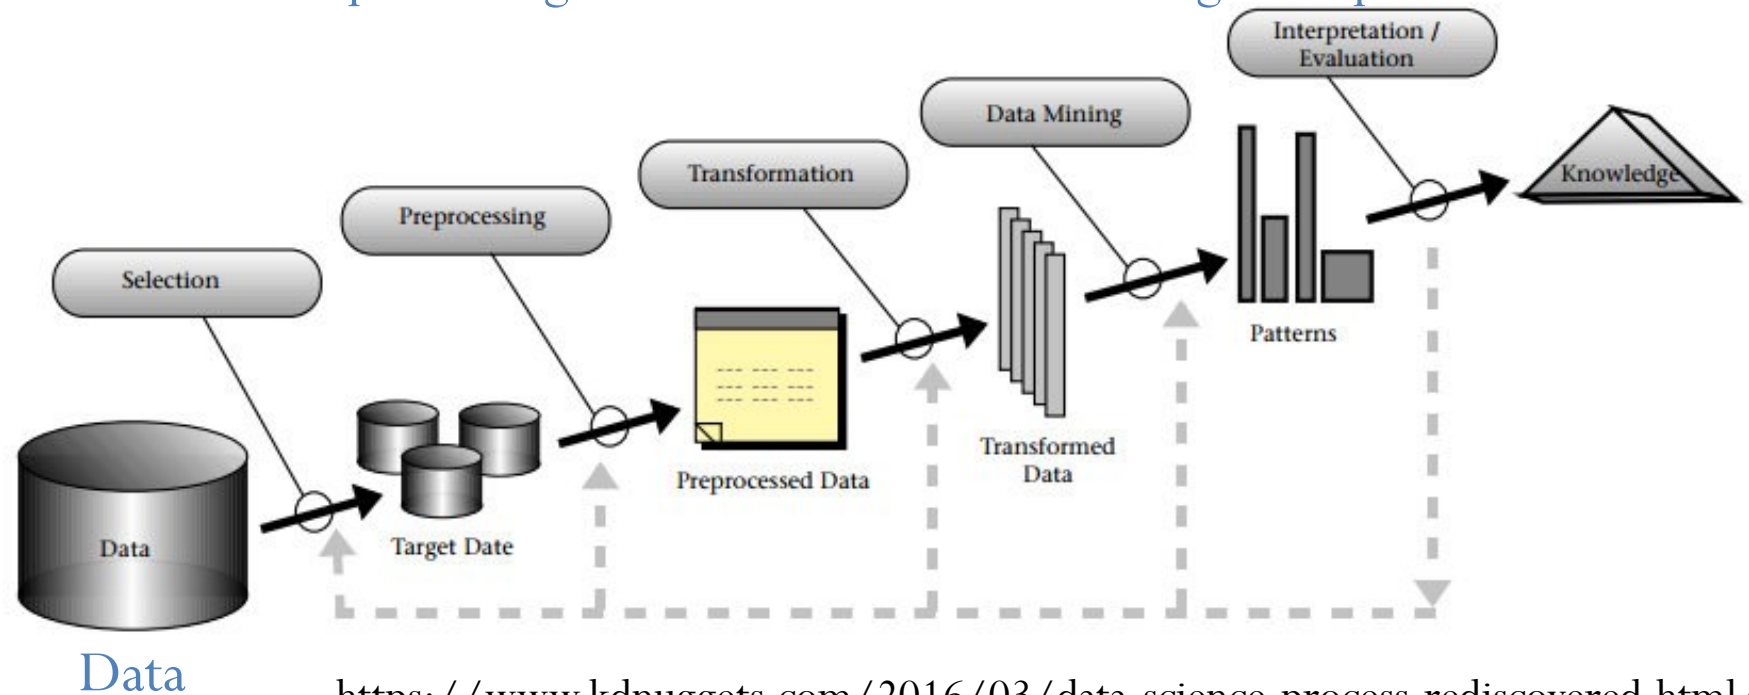
Fayyad, Usama M., Gregory Piatetsky-Shapiro, and Padhraic Smyth.

"Knowledge Discovery and Data Mining: Towards a Unifying Framework." *KDD*. 1996.

Knowledge Discovery and Data Mining

- “a unifying framework for Knowledge Discovery in Database (KDD)”
- links between data mining, knowledge discovery, and other related field

Selection, Preprocessing, Transformation, Data Mining, Interpretation/Evaluation



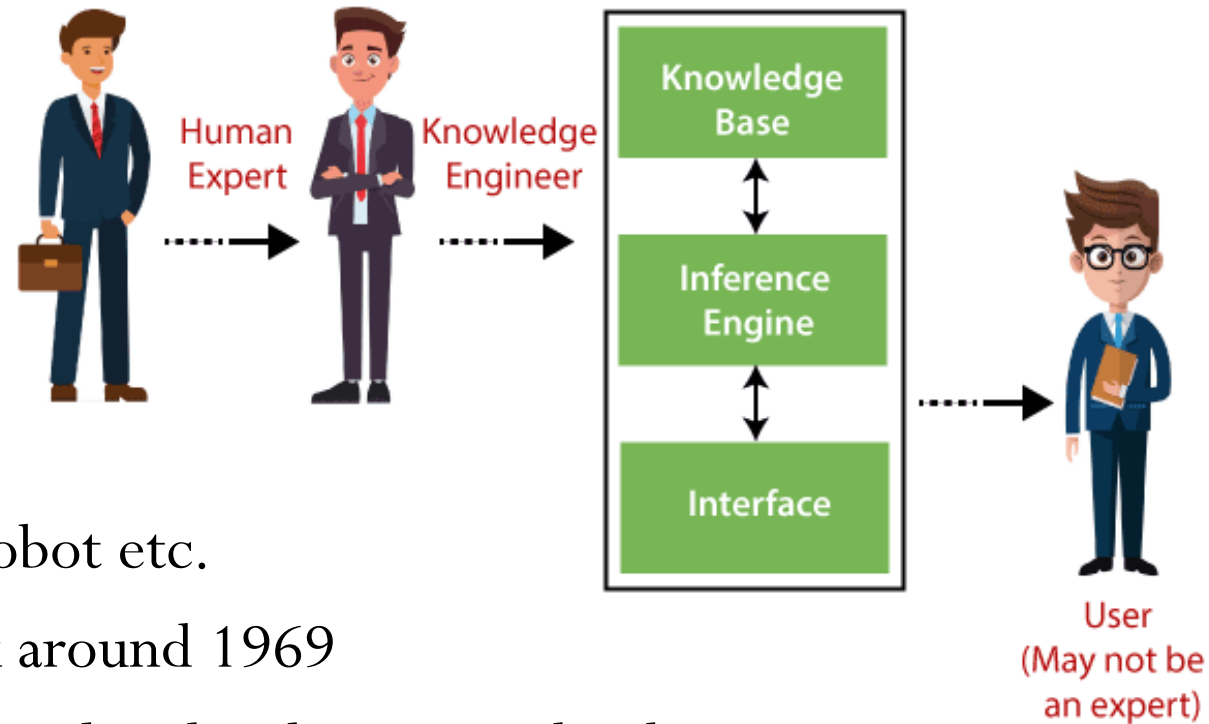
<https://www.kdnuggets.com/2016/03/data-science-process-rediscovered.html/2>

Fayyad, Usama M., Gregory Piatetsky-Shapiro, and Padhraic Smyth.

"Knowledge Discovery and Data Mining: Towards a Unifying Framework." **KDD**. 1996.

Expert, System, and Knowledge Engineer

- Expert Systems (1960-74)
 - Explicit rules
- Playing Chess,
- Organic and Biology recommendations
- Solving word problem in Algebra
- Natural Language processing, Mobile Robot etc.
- Backpropagation based Neural Network around 1969
- Association Rule Mining by Rakesh Agrawal and Srikant Ramakrishnan 1993-95
- Big Files and Google File System in 1995-2005 by Larry Page, Sergey Brin and Sanjay Ghemawat, et al.




Big Files to Google File System (GFS)

- Earlier Google effort, "BigFiles", developed by Larry Page and Sergey Brin.
 - Supervisors: Hector Garcia-Molina, Rajeev Motwani, Jeff Ullman, and Terry Winograd
- "Big File" was regenerated as "Google File System" by Sanjay Ghemawat, et al.
- Google File System (GFS)
 - "It is widely deployed within Google as the storage platform for the generation and processing of data used by Google service as well as research and development efforts that require large data sets." 2003
 - "The largest cluster to date provides hundreds of terabytes of storage across thousands of disks on over a thousand machines, and it is concurrently accessed by hundreds of clients." 2003

<http://infolab.stanford.edu/~backrub/google.html>

Sanjay Ghemawat, Howard Gobioff, and Shun-Tak Leung. "The Google file system." Proceedings of the nineteenth ACM symposium on Operating systems principles. 2003.

Google File System (GFS)



https://en.wikipedia.org/wiki/Google_File_System

<https://sites.google.com/site/gfsassignmentwiki/home>

Sanjay Ghemawat, Howard Gobioff, and Shun-Tak Leung. "The Google file system." Proceedings of the nineteenth ACM symposium on Operating systems principles. 2003.

GFS to HDFS

- Google File System (GFS) has similar open-source
 - “Hadoop Distributed File System (HDFS)”
- GFS and HDFS are distributed computing environment to process “Big Data”.
- GFS and HDFS are not implemented in the kernel of an operating system, but they are instead provided as a userspace library.
- GFS and HDFS properties
 - Files are divided into fixed-size chunks of 64 megabytes.
 - Scalable distributed file system for large distributed data intensive applications.
 - Provides fault tolerance.
 - High aggregate performance to a large number of clients.

https://en.wikipedia.org/wiki/Google_File_System

Sanjay Ghemawat, Howard Gobioff, and Shun-Tak Leung. "The Google file system." Proceedings of the nineteenth ACM symposium on Operating systems principles. 2003.

Big Data

- Big data can be described by the following characteristics:
 - Volume: size large than terabytes and petabytes
 - Variety: type and nature, structured, semi-structured or unstructured
 - Velocity: speed of generation and processing to meet the demands
 - Veracity: the data quality and the data value
 - Value: Useful or not useful
- The main components and ecosystem of Big Data
 - Data Analytics: data mining, machine learning and natural language processing etc.
 - Technologies: Business Intelligence, Cloud computing & Databases etc.
 - Visualization: Charts, Graphs etc.



Hadoop to handle Big Data


- “The Apache Hadoop project develops open source software for reliable, scalable, distributed computing.” Software library and a framework.
- Created by Doug Cutting Named on his son's stuffed elephant
- For distributed processing of large data sets across clusters of computers using simple programming models.
- Locality of reference
- Scalability: Scale up from single servers to thousands of machines,
 - Each offering local computation and storage
 - Program remains same for 10, 100, 1000,... nodes
 - Corresponding performance improvement
- Fault-tolerant file system:
 - Detect and handle failures and Delivering a highly-available.
- Hadoop Distributed File System (HDFS) Modeled on Google File system
- MapReduce for Parallel computation using
- Components – Pig, Hbase, HIVE, ZooKeeper



Apache Spark



- Unified analytics engine for large-scale data processing.
- Speed: Run workloads 100x faster.
- Both batch and streaming data, using Directed Acyclic Graph (DAG) scheduler, a query optimizer, and a physical execution engine.
- Ease of Use: Write applications quickly in Java, Scala, Python, R, and SQL.
- Spark offers 80+ high-level operators to build parallel apps.



Spark: Unified Big Data Analytics

- New applications of Big data workloads on Unified Engine of
 - Streaming, Batch, and Interactive.
- Composability in programming libraries for big data and encourages development of interoperable libraries
- Combining the SQL, machine learning, and streaming libraries in Spark

```
// Load historical data as an RDD using Spark SQL
val trainingData = sql(
  "SELECT location, language FROM old_tweets")

// Train a K-means model using MLlib
val model = new KMeans()
  .setFeaturesCol("location")
  .setPredictionCol("language")
  .fit(trainingData)
// Apply the model to new tweets in a stream
TwitterUtils.createStream(...)
  .map(tweet => model.predict(tweet.location))
```

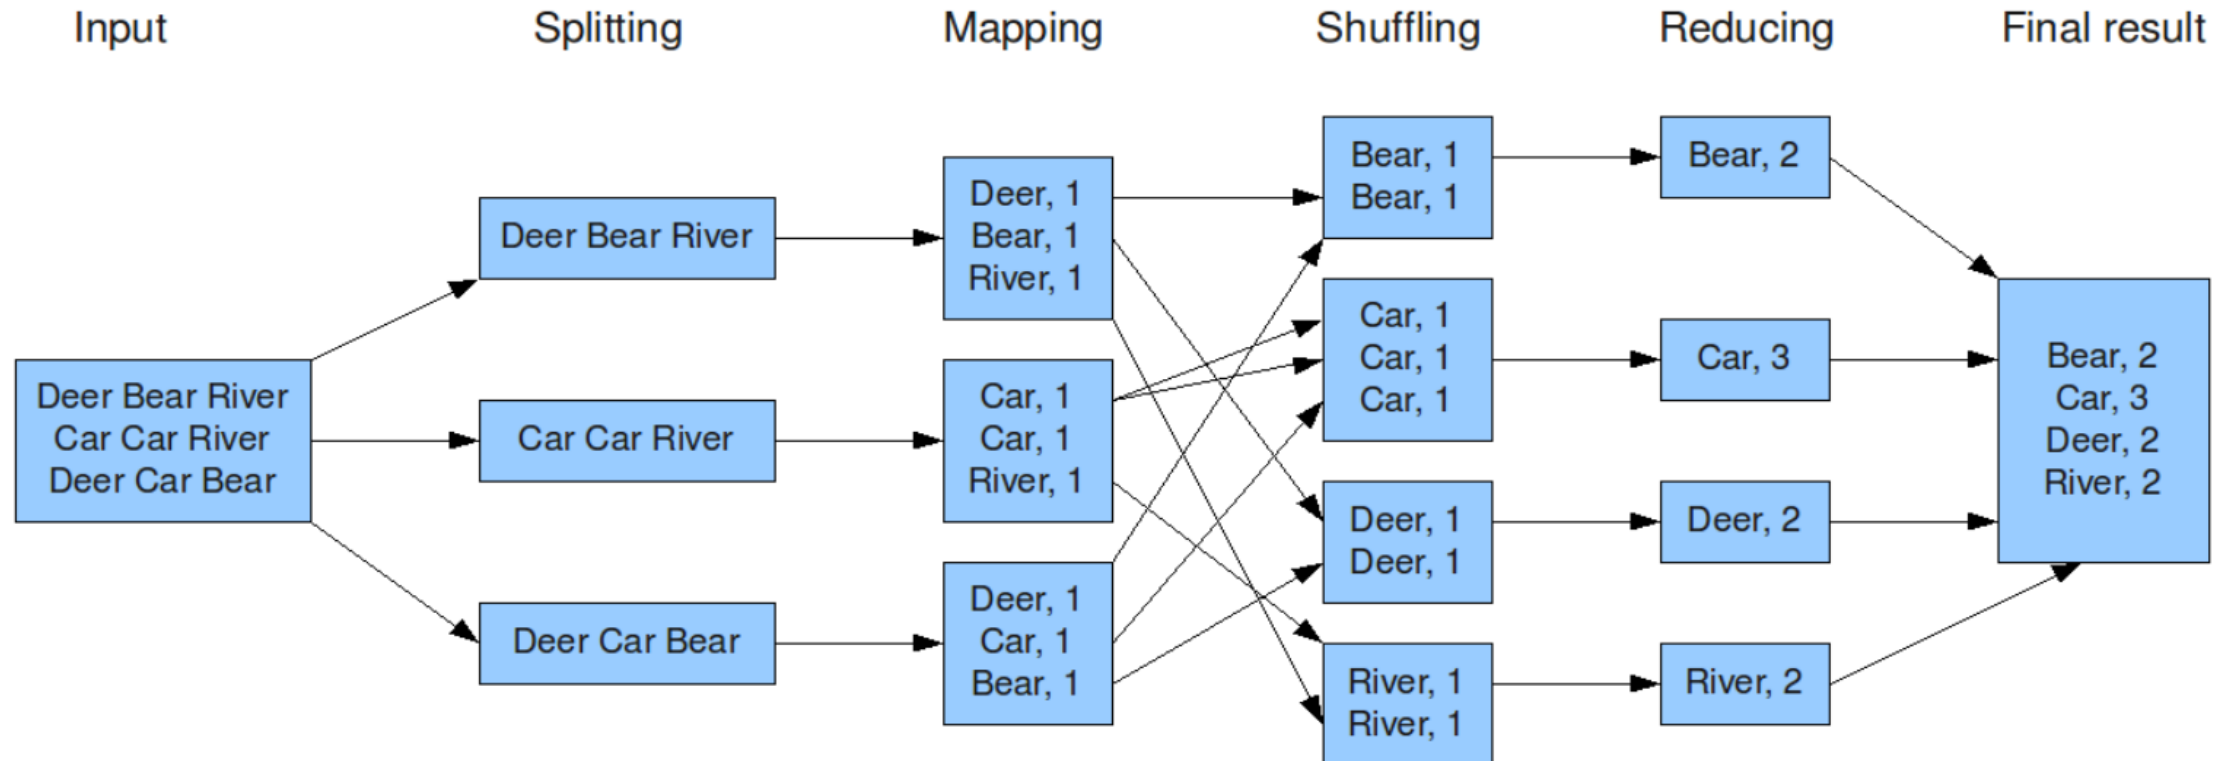
Apache Spark



- Combine SQL, streaming, and complex analytics. Spark libraries
 - [SQL and DataFrames](#),
 - [MLlib](#) for machine learning,
 - [GraphX](#), and
 - [Spark Streaming](#).
- Spark runs on Hadoop, Apache Mesos, Kubernetes, standalone, or in the cloud.
- Run Spark using its [standalone cluster mode](#), on [EC2](#), on [Hadoop YARN](#), on [Mesos](#), or on [Kubernetes](#).
- Access data in [HDFS](#), [Alluxio](#), [Apache Cassandra](#), [Apache HBase](#), [Apache Hive](#), and hundreds of other data sources.

Word Count

- **Map:** Input lines of text to breaks them into words gives outputs for each word
<key = word, value =1 >
- **Reduce:** Input <word, 1> output <word, + value>



Spark Code Example



- Word Count


```
JavaRDD<String> textFile = sc.textFile("hdfs://...");
JavaPairRDD<String, Integer> counts = textFile
    .flatMap(s -> Arrays.asList(s.split(" ")).iterator())
    .mapToPair(word -> new Tuple2<>(word, 1))
    .reduceByKey((a, b) -> a + b);
counts.saveAsTextFile("hdfs://...");
```

- Pi Estimation

```
List<Integer> l = new ArrayList<>(NUM_SAMPLES);
for (int i = 0; i < NUM_SAMPLES; i++) {
    l.add(i);
}
```

```
long count = sc.parallelize(l).filter(i -> {
    double x = Math.random();
    double y = Math.random();
    return x*x + y*y < 1;
}).count();
System.out.println("Pi is roughly " + 4.0 * count / NUM_SAMPLES);
```

$$\text{Pi} = 4 \times \frac{\text{number of random point inside the circle}}{\text{number of random point inside the square}}$$



Spark Streaming



- Build scalable fault-tolerant streaming applications.
- Write streaming jobs -- Same Way -- Write batch jobs

- Counting tweets on a sliding window

```
TwitterUtils.createStream(...)  
  .filter(_.getText.contains("Spark"))  
  .countByWindow(Seconds(5))
```

- Reuse the same code for batch processing

- Find words with higher frequency than historic data:

```
stream.join(historicCounts).filter {  
  case (word, (curCount, oldCount)) =>  
    curCount > oldCount  
}
```

Batch
processing
takes N unit
time to
process M
unit of data

Batch
processing
takes N+x
unit time
to process
M+y unit
of data

Stream
processing
takes N unit
time to
process M
unit of data

Stream
processing
takes x
unit time
to process
M+y unit
of data



Spark MLlib

- Spark's scalable machine learning library
- Spark MLlib algorithms
 - Classification: logistic regression, naive Bayes,...
 - Regression: generalized linear regression, survival regression,...
 - Decision trees, Random forests, and Gradient-boosted trees
 - Recommendation: Alternating Least Squares (ALS)
 - Clustering: K-means, Gaussian mixtures (GMMs),...
 - Topic modeling: Latent Dirichlet Allocation (LDA)
 - **Frequent itemsets, Association rules, and Sequential pattern mining**



Spark MLlib

- Spark's scalable machine learning library
- Spark MLlib algorithms
 - Classification: logistic regression, naive Bayes,...
 - Regression: generalized linear regression, survival regression,...
 - Decision trees, Random forests, and Gradient-boosted trees
 - Recommendation: Alternating Least Squares (ALS)
 - Clustering: K-means, Gaussian mixtures (GMMs),...
 - Topic modeling: Latent Dirichlet Allocation (LDA)
 - **Frequent itemsets, Association rules, and Sequential pattern mining**

Apriori algorithm: Association Rule Mining

- Let $I = \{i_1, i_2, \dots, i_m\}$ be a set of literals, called items.
- *Support* of a rule $X \rightarrow Y$ is the percentage of transactions that contain both X and Y .
- *Confidence* of a rule is percentage the if-then statements ($X \rightarrow Y$) are found true
- Find all rules that satisfy a user-specified *minimum support* and *minimum confidence*

TID	Transaction Items
1	Bread, Jelly, PeanutButter
2	Bread, PeanutButter
3	Bread, Milk, PeanutButter
4	Beer, Bread
5	Beer, Milk



$\{\text{Bread}\} \rightarrow \{\text{PeanutButter}\}$ (Sup = 60%, Conf = 75%)
 $\{\text{PeanutButter}\} \rightarrow \{\text{Bread}\}$ (Sup = 60%, Conf = 100%)
 $\{\text{Beer}\} \rightarrow \{\text{Bread}\}$ (Sup = 20%, Conf = 50%)
 $\{\text{PeanutButter}\} \rightarrow \{\text{Jelly}\}$ (Sup = 20%, Conf = 33.33%)
 $\{\text{Jelly}\} \rightarrow \{\text{PeanutButter}\}$ (Sup = 20%, Conf = 100%)
 $\{\text{Jelly}\} \rightarrow \{\text{Milk}\}$ (Sup = 0%, Conf = 0%)


Apriori algorithm: Association Rule Mining

- Let $I = \{i_1, i_2, \dots, i_m\}$ be a set of literals, called items.
- *Support* of a rule $X \rightarrow Y$ is the percentage of transactions that contain both X and Y .
- *Confidence* of a rule is percentage the if-then statements ($X \rightarrow Y$) are found true
- Find all rules that satisfy a user-specified *minimum support* and *minimum confidence*
 - 100% of transactions that purchase *PeanutButter* (antecedent) also purchase *Bread* (consequent). The number 100% is the confidence factor of the rule
 - $[PeanutButter] \rightarrow [Bread]$ (Sup = 60%, Conf = 100%)
 - 98% of customers who purchase *Tires* and *Auto accessories* also buy some *Automotive services*; here 98% is called the confidence of the rule.
 - $[Auto Accessories], [Tires] \rightarrow [Automotive Services]$ 98%

FP-Growth for recommendation

- “FP” stands for Frequent Pattern in a Dataset of transactions
 1. calculate item frequencies and identify frequent items,
 2. a suffix tree (FP-tree) structure to encode transactions, and
 3. frequent itemsets can be extracted from the FP-tree.
- Input: Transaction database
- Intermediate Output: FP-Tree
- Output: $\{f, c, a \rightarrow a, m, p\}, \{f, c, a \rightarrow b, m\}$

TID	Items Bought	(Ordered) Frequent Items
100	<i>f, a, c, d, g, i, m, p</i>	<i>f, c, a, m, p</i>
200	<i>a, b, c, f, l, m, o</i>	<i>f, c, a, b, m</i>
300	<i>b, f, h, j, o</i>	<i>f, b</i>
400	<i>b, c, k, s, p</i>	<i>c, b, p</i>
500	<i>a, f, c, e, l, p, m, n</i>	<i>f, c, a, m, p</i>



PFP: Parallel FP-Growth

- In Spark ML-Library (MLLib), a parallel version of FP-growth called PFP: Parallel FP-Growth
- PFP distributes the work of growing FP-trees based on the suffixes of transactions.
- More scalable than a single-machine implementation.
- PFP partitions computation, where each machine executes an independent group of mining tasks

PFP: Parallel FP-Growth

Example of MapReduce FP-Growth: Five transactions composed of lower-case alphabets representing items

Map inputs (transactions) key="": value	Sorted transactions (with infrequent items eliminated)	Map outputs (conditional transactions) key: value	Reduce inputs (conditional databases) key: value	Conditional FP-trees
f a c d g i m p	f c a m p	p: f c a m m: f c a a: f c c: f	p: { f c a m / f c a m / c b }	{(c:3)} p
a b c f l m o	f c a b m	m: f c a b b: f c a a: f c c: f	m: { f c a / f c a / f c a b }	{ (f:3, c:3, a:3) } m
b f h j o	f b	b: f	b: { f c a / f / c }	{ } b
b c k s p	c b p	p: c b b: c	a: { f c / f c / f c }	{ (f:3, c:3) } a
a f c e l p m n	f c a m p	p: f c a m m: f c a a: f c c: f	c: { f / f / f }	{ (f:3) } c

Li, Haoyuan, et al. "PFP: Parallel FP-Growth for query recommendation."

Proceedings of the 2008 ACM Conference on Recommender systems. 2008.

PFP: Parallel FP-Growth

- FP-Growth implementation takes the following (hyper-)parameters
 - minSupport: the minimum support for an itemset to be identified as frequent e.g., if an item appears 3 out of 6 transactions, it has a support of $3/6=0.5$.
 - minConfidence: minimum confidence for generating Association Rule e.g., if in the transactions itemset X appears 5 times, X and Y co-occur only 3 times, the confidence for the rule $X \Rightarrow Y$ is then $3/5 = 0.6$.
 - numPartitions: the number of partitions used to distribute the work.
- FP-Growth model provides:
 - freqItemsets: frequent itemsets in the format of DataFrame(“items”[Array], “freq”[Long])
 - associationRules: association rules generated with confidence above minConfidence, in the format of DataFrame(“antecedent”[Array], “consequent”[Array], “confidence”[Double]).

PFP: Parallel FP-Growth

```
List<Row> data = Arrays.asList(
    RowFactory.create(Arrays.asList("1 2 5".split(" "))),
    RowFactory.create(Arrays.asList("1 2 3 5".split(" "))),
    RowFactory.create(Arrays.asList("1 2".split(" ")))
);

StructType schema = new StructType(new StructField[]{ new StructField(
    "items", new ArrayType(DataTypes.StringType, true), false, Metadata.empty())
});

Dataset<Row> itemsSDF = spark.createDataFrame(data, schema);

FPGrowthModel model = new FPGrowth()
    .setItemsCol("items")
    .setMinSupport(0.5)
    .setMinConfidence(0.6)
    .fit(itemsSDF);

// Display frequent itemsets.
model.freqItemsets().show();

// Display generated association rules.
model.associationRules().show();

// transform examines the input items against all the association rules and summarize the
// consequents as prediction
model.transform(itemsSDF).show();
```

<https://spark.apache.org/docs/3.3.1/ml-frequent-pattern-mining.html>

תודה רבה

Hebrew

Ευχαριστώ

Greek

Спасибо

Russian

Danke

German

Merci

French

धन्यवादः

Sanskrit

நன்றி

Tamil

شكراً

Arabic

ಧನ್ಯವಾದಗಳು

Kannada

Thank You

English

നന്നി

Malayalam

Grazie

Italian

ధన్యవాదాలు

Telugu

આભાર

Gujarati

多謝

Traditional Chinese

Gracias

Spanish

ਧੰਨਵਾਦ

Punjabi

धन्यवाद

Hindi & Marathi

多谢

Simplified Chinese

<https://sites.google.com/site/animeshchaturvedi07>

Obrigado

Portuguese

ありがとうございました

Japanese

ขอบคุณ

Thai

감사합니다

Korean